

K NEAREST NEIGHBORS

This PDF version is not interactive. For interactivity, use the [html version instead](#)

Carsten Lange
clange@cpp.edu
Cal Poly, Pomona

BEFORE WE BEGIN LET US DO A THOUGHT EXPERIMENT

I need three volunteers. Everybody else, please follow along on one computer in groups of three.

Imagine the three of you are moving into a shared apartment.

There are two bedrooms: two people will share a big room, and one person will have a small room alone.

To minimize conflicts, we want the two most similar people to share the two rooms.

Please rate yourselves on a **scale from 1 to 10** for the two criteria below:

1. How tidy you are (**1=messy, 10= tidy**)
2. How social you are at home (**1=quiet, 10=very social**)

We will treat your answers as data points and use Euclidean distance to decide who should share.

THOUGHT EXPERIMENT: DATA ENTRY

For `Student1`, `Student2`, and `Student3`, let us enter the names of the three volunteer student. Then we replace the sample scores (e.g., `Tidy = c(1, 8, 3)`) with the scores from the volunteer students (in the correct order of `Student1`, `Student2`, and `Student3`).

Afterwards, click **Run Code** (note, you need to run this code in order to work with the following examples).

R Code  Start Over

 Run Code

```
1 library(kableExtra)
2 DataRoomMates <- data.frame(
3   name = c("Student 1", "Student 2", "Student 3"),
4   Tidy = c(1, 8, 3),    # 1 = messy, 10 = very tidy
5   Social = c(7, 2, 5)    # 1 = quiet, 10 = very social
6 )
7 kable(DataRoomMates)
```

THOUGHT EXPERIMENT: PLOTTING THE DATA

R Code [⟳ Start Over](#)

▷ Run Code

```
1 library(ggplot2)
2 ggplot(DataRoomMates, aes(x = Tidy, y = Social, label = name)) +
3   geom_point(size = 4) +
4   geom_text(vjust = -1) +
5   scale_x_continuous(limits = c(0.5, 10.5), breaks=seq(1:10)) +
6   scale_y_continuous(limits = c(0.5, 10.5), breaks=seq(1:10)) +
7   coord_fixed(ratio = 1) +    # <-- makes the plot square
8   labs(
9     title = "Roommate Matching Using Euclidean Distance",
10    x = "Tidy (1 = messy, 10 = tidy)",
11    y = "Social at Home (1 = quiet, 10 = social)"
12  )
```

$$EucDist^2 = DiffTidy^2 + DiffSocial^2$$

$$EucDist = \sqrt{DiffTidy^2 + DiffSocial^2}$$

COMPUTE EUCLIDEAN DISTANCES

R Code

⟳ Start Over

▷ Run Code

```
1 MatrixDistance <- dist(DataRoomMates[, c("Tidy", "Social")], method = "euclidean")
2 print(as.matrix(MatrixDistance))
```

Testing the distance between the first and the second student:

$$EucDist^2 = DiffTidy^2 + DiffSocial^2$$

$$EucDist = \sqrt{DiffTidy^2 + DiffSocial^2}$$

R Code

⟳ Start Over

▷ Run Code

```
1 DiffTidy=DataRoomMates[[1,2]]-DataRoomMates[[2,2]]
2 cat("The difference in Tidy between the first and the second student is:",DiffTidy)
3 DiffSocial=DataRoomMates[[1,3]]-DataRoomMates[[2,3]]
4 cat("The difference in Social between the first and the second student is:",DiffSocial)
5 EucDist=sqrt((DiffTidy^2+DiffSocial^2))
6 cat("Euclidean Distance=",EucDist)
```

USE THE SAME SCALE FOR ALL PREDICTOR VARIABLES

IT IS VERY IMPORTANT THAT ALL PREDICTOR VARIABLES HAVE THE SAME SCALE

Like here, both variables are in a range from 1 to 10.

WHAT HAPPENS IF WE SCALE **Social** FROM 10 TO 100 INSTEAD FROM 1 TO 10

R Code ⟳ Start Over

▷ Run Code

```
1 DataRoomMatesScaled>DataRoomMates |>  
2   mutate(Social=Social*10)  
3  
4 kable(DataRoomMatesScaled)
```

THE INFLUENCE OF SOCIAL COMPARED TO TIDY BECOMES VERY SMALL

R Code

⟳ Start Over

▷ Run Code

```
1 library(ggplot2)
2 ggplot(DataRoomMatesScaled, aes(x = Tidy, y = Social, label = name)) +
3   geom_point(size = 4) +
4   geom_text(vjust = -1) +
5   scale_x_continuous(limits = c(1, 10), breaks=seq(1,10,9)) +
6   scale_y_continuous(limits = c(10, 100), breaks=seq(10,100,10)) +
7   coord_fixed(ratio = 1) +    # <-- makes the plot square
8   labs(
9     title = "Roommate Matching Using Euclidean Distance",
10    x = "Tidy (1 = messy, 10 = tidy)",
11    y = "Social at Home (1 = quiet, 10 = social)"
12  )
```

The Euclidean Distance is almost the same as `DistSocial`, which makes `DistTidy` almost irrelevant.

OVERVIEW

In this session you will learn:

1. What is the underlying **idea of k-Nearest Neighbors**
2. How similarity can be measured with **Euclidean distance**
3. Why **scaling predictor variables** is important for some machine learning models
4. Why the **tidymodels package** makes it easy to work with machine learning models
5. How you can define a **recipe** to pre-process data with the **tidymodels** package
6. How you can define a **model-design** with the **tidymodels** package
7. How you can create a machine learning **workflow** with the **tidymodels** package
8. How **metrics** derived from a **confusion matrix** can be used to asses prediction quality
9. Why you have to be careful when interpreting *accuracy*, when you work with **unbalanced observations**
10. How a machine learning model can **process images** and how OCR (Optical Character Recognition) works

ABOUT THE WINE DATASET

We will work with a publicly available wine dataset¹ containing 3,198 observations about different wines and their chemical properties.

Our goal is to develop a k-Nearest Neighbors model that can predict if a wine is red or white based on the wine's chemical properties.

1. Cortez, Paulo, António Cerdeira, Fernando Almeida, Telmo Matos, and José Reis. 2009. "Modeling Wine Preferences by Data Mining from Physicochemical Properties." *Decision Support Systems* 47 (4): 547–53. <https://doi.org/10.1016/j.dss.2009.05.016>.

RAW OBSERVATIONS FROM WINE DATASET

```
1 library(rio)
2 DataWine=import("https://ai.lange-analytics.com/data/WineData.rds")
3 print(DataWine)
```

	wineColor	acidity	volatileAcidity	citricAcid	residualSugar	Chlorides
1	red	10.80	0.320	0.44	1.60	0.063
2	white	6.40	0.310	0.39	7.50	0.040
3	white	9.40	0.280	0.30	1.60	0.045
4	white	8.20	0.220	0.36	6.80	0.034
5	white	6.40	0.290	0.44	3.60	0.197
6	red	6.70	0.855	0.02	1.90	0.064
7	red	11.80	0.380	0.55	2.10	0.071
8	white	6.70	0.250	0.23	7.20	0.038
9	red	7.50	0.380	0.57	2.30	0.106
10	red	7.10	0.270	0.60	2.10	0.074
11	white	6.40	0.270	0.19	1.90	0.085
12	red	7.80	0.600	0.26	2.00	0.080
13	red	8.00	0.580	0.28	3.20	0.066
14	white	7.00	0.360	0.35	2.50	0.048
15	red	9.90	0.440	0.46	2.20	0.091
16	white	7.80	0.280	0.31	2.10	0.046
17	red	7.60	0.400	0.29	1.90	0.078
18	white	5.90	0.260	0.24	2.40	0.046
19	white	6.80	0.180	0.28	9.80	0.039
20	white	6.80	0.250	0.30	11.80	0.043
21	white	6.90	0.190	0.35	1.70	0.036
22	white	6.50	0.310	0.14	7.50	0.044

OBSERVATIONS FROM WINE DATASET FOR SELECTED VARIABLES

SULFUR DIOXIDE AND ACIDITY

Note we use `clean_names("upper_camel")` from the `janitor` package to change all column (variable) names to UpperCamel.

```
1 library(tidyverse); library(rio);library(janitor)
2 DataWine=import("https://ai.lange-analytics.com/data/WineData.rds") |>
3   clean_names("upper_camel") |>
4   select(WineColor,Sulfur=TotalSulfurDioxide,Acidity) |>
5   mutate(WineColor=as.factor(WineColor))
6 print(DataWine)
```

	WineColor	Sulfur	Acidity
1	red	37.0	10.80
2	white	213.0	6.40
3	white	139.0	9.40
4	white	90.0	8.20
5	white	183.0	6.40
6	red	38.0	6.70
7	red	19.0	11.80
8	white	220.0	6.70
9	red	12.0	7.50
10	red	25.0	7.10
11	white	196.0	6.40
12	red	131.0	7.80
13	red	114.0	8.00
14	white	161.0	7.00
15	red	41.0	9.90
16	white	208.0	7.80
17	red	66.0	7.60
18	white	132.0	5.90
19	white	113.0	6.80
20	white	133.0	6.80
21	white	101.0	6.90

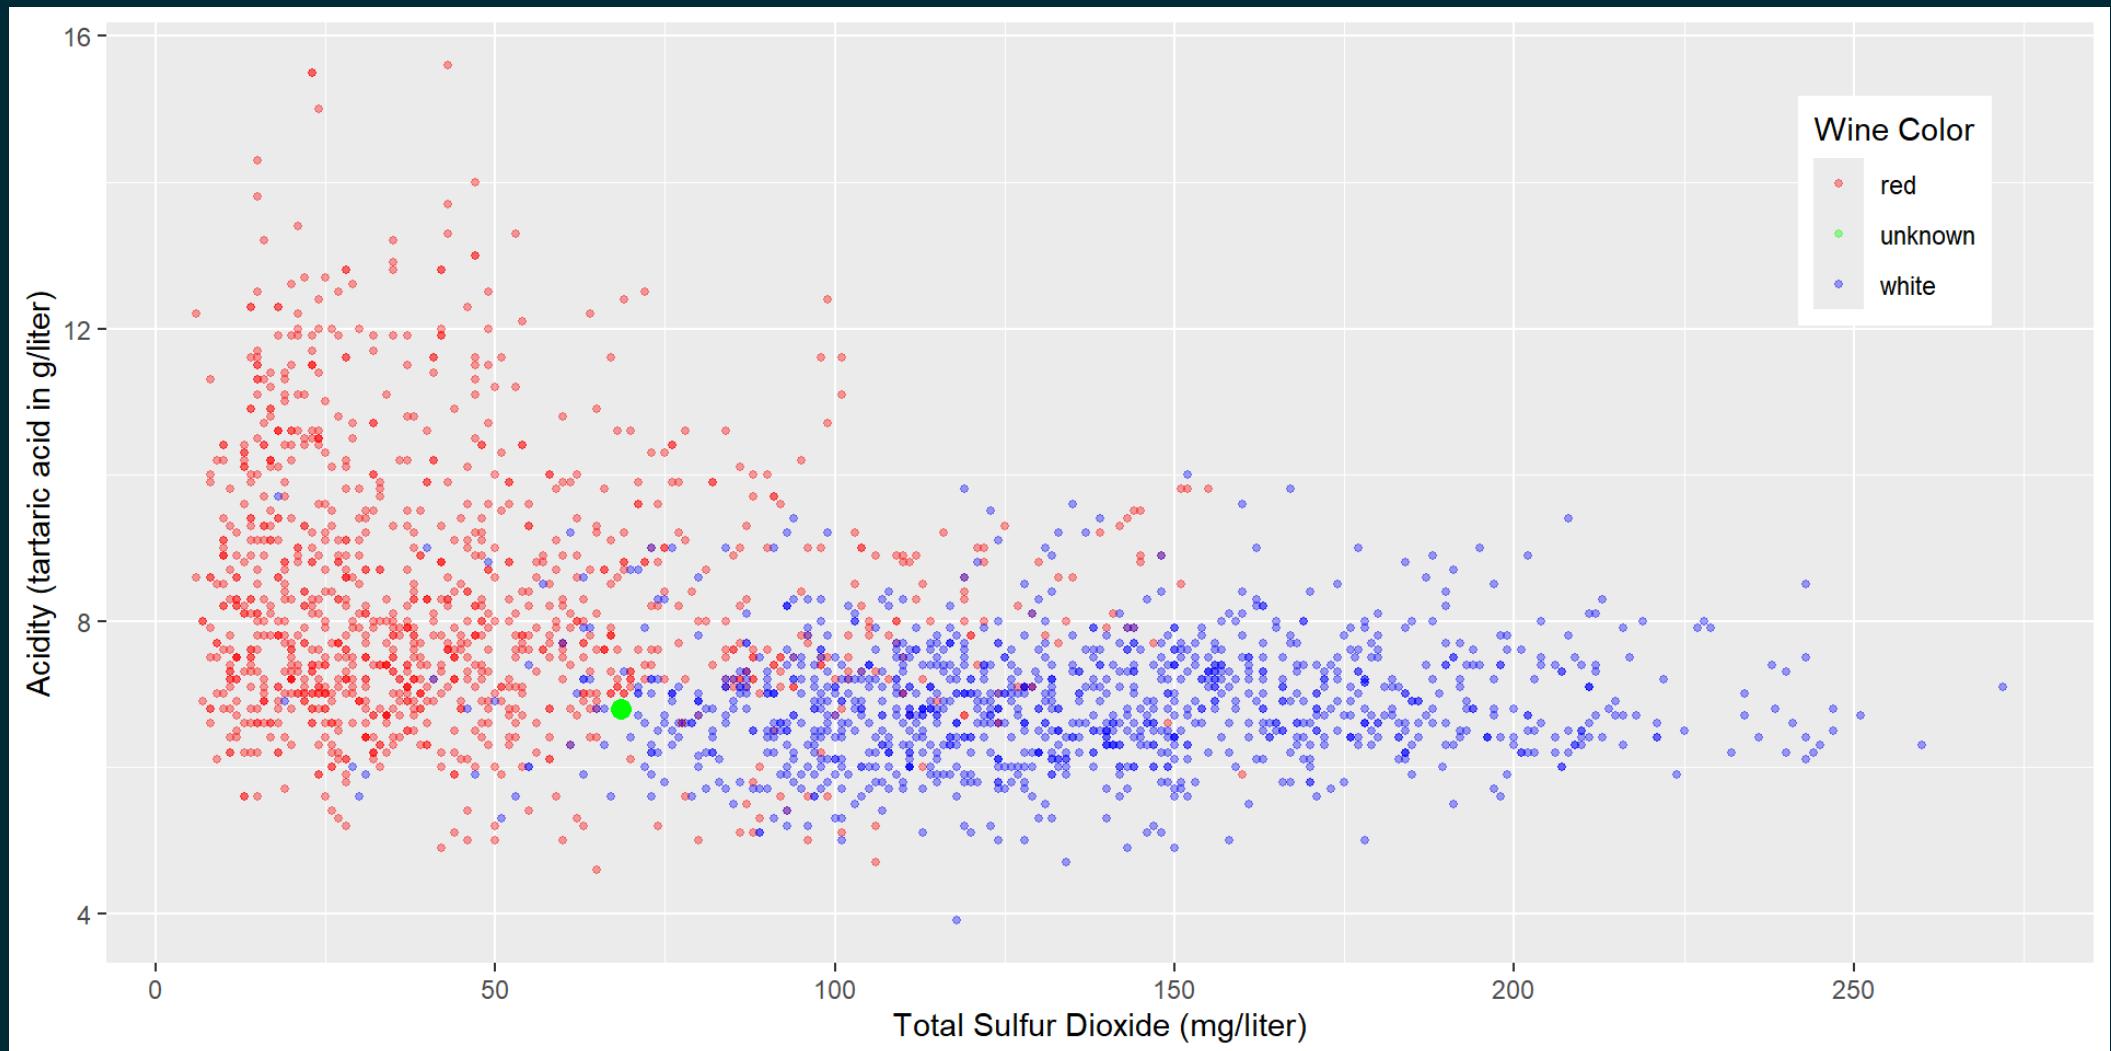
CONCEPTUAL DETOUR: BEFORE STARTING WITH K NEAREST NEIGHBORS

INTRODUCING A FEW OTHER MACHINE LEARNING MODELS WITH DIAGRAMS

LET US FIND SOME EYEBALLING TECHNIQUES THAT ARE RELATED TO VARIOUS MACHINE LEARNING MODELS

EYE BALLING TECHNIQUES TO IDENTIFY RED AND WHITE WINES

TRY EYEBALLING THE DATA

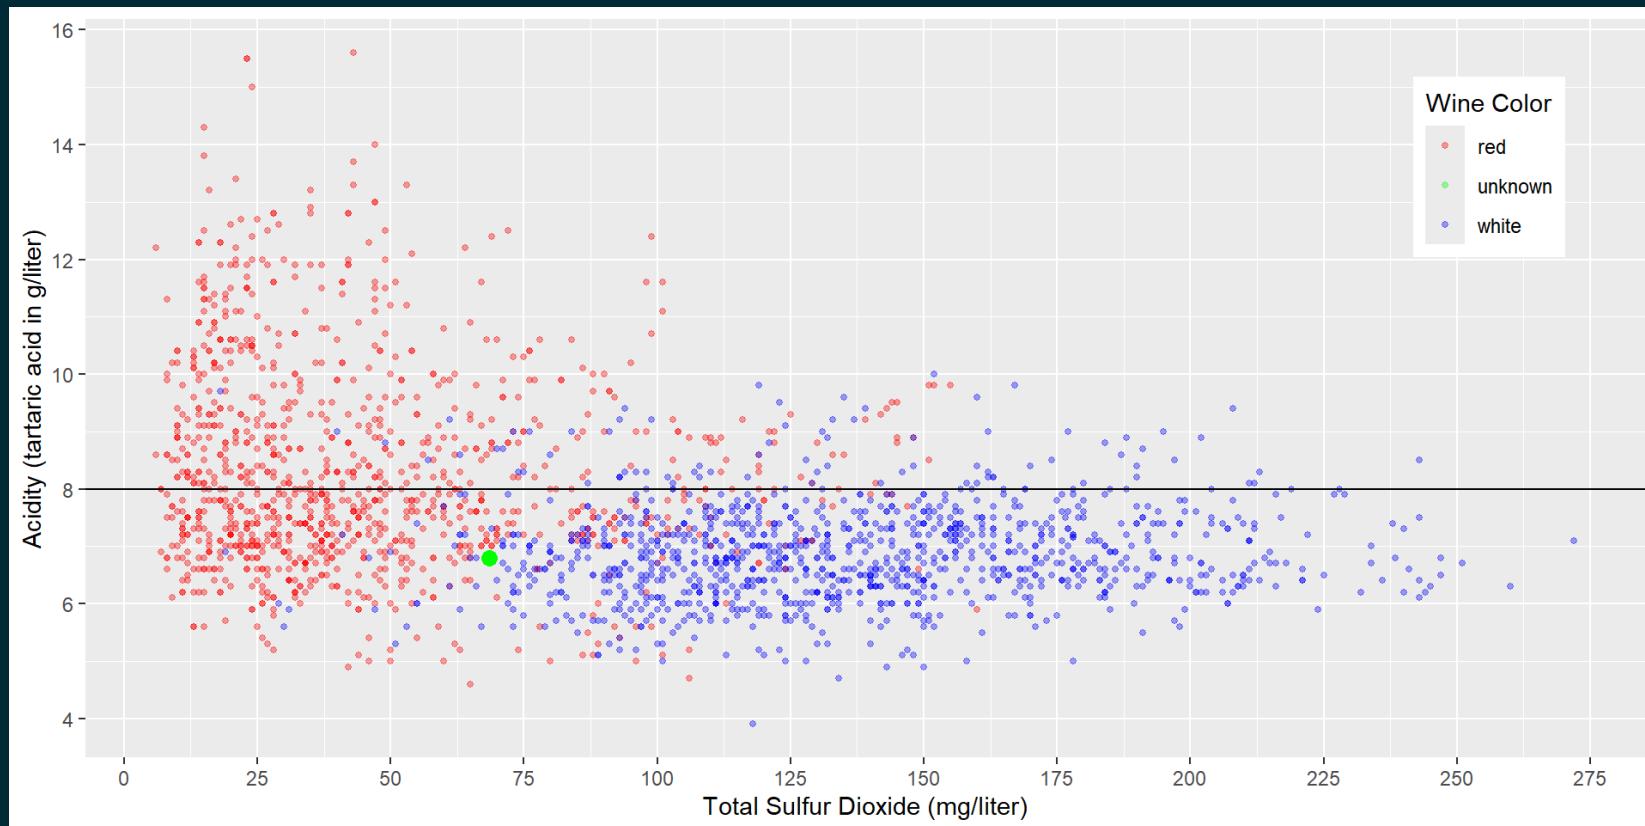


Acidity and Total Sulfur Dioxide Related to Wine Color

Textbook

EYE BALLING TECHNIQUES TO IDENTIFY RED AND WHITE WINES

HORIZONTAL BOUNDARY



Horizontal Decision Boundary for Acidity and Total Sulfur Dioxide Related to Wine Color

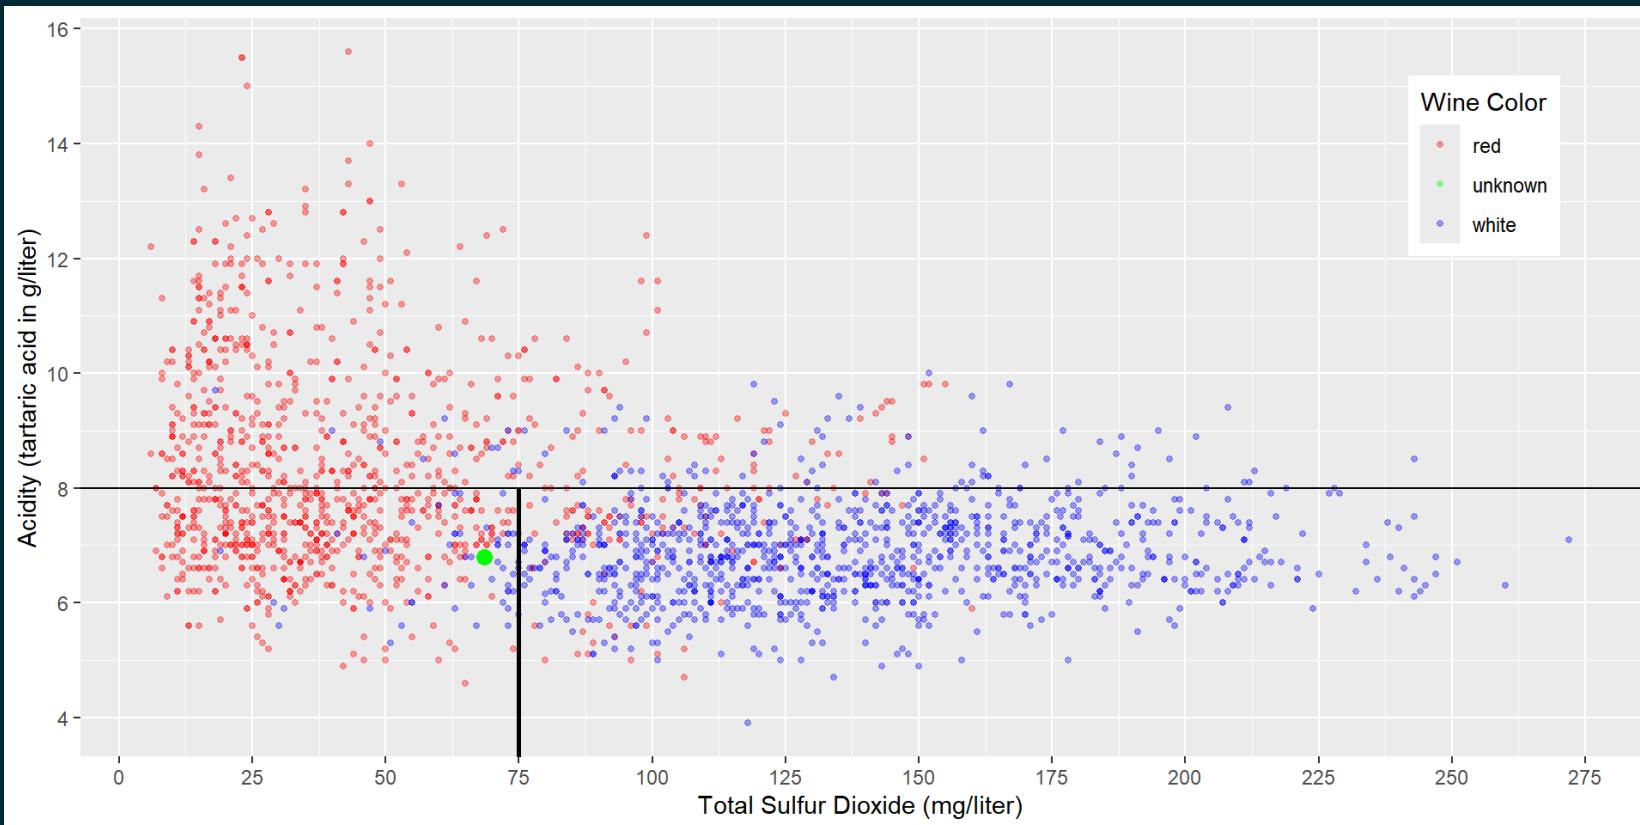
CONFUSION MATRIX

		Truth	
		Red Wine	White Wine
Prediction	Red Wine	TP: 'half'	FP: 'few'
	White Wine	FN: 'half'	TN: 'most'

Textbook

EYEBALLING TECHNIQUES TO IDENTIFY RED AND WHITE WINES

CREATING SUBSPACES LIKE SIMILAR TO A DECISION TREE



Sub-Space Boundaries for Acidity and Total Sulfur Dioxide Related to Wine Color

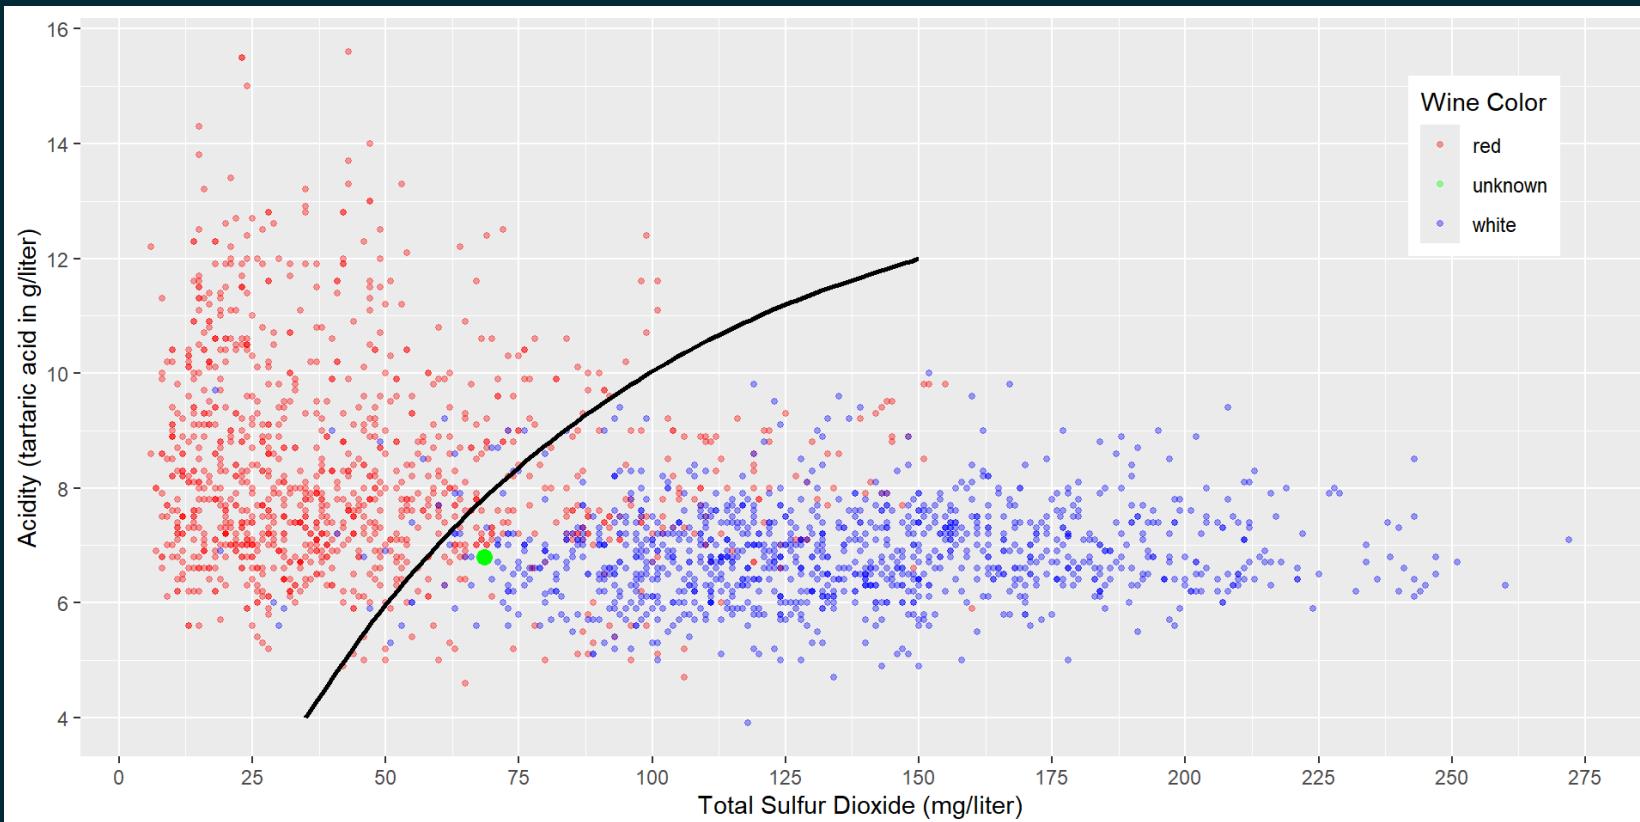
CONFUSION MATRIX

		Truth	
		Red Wine	White Wine
Prediction	Red Wine	TP: 'most'	FP: 'few'
	White Wine	FN: 'few'	TN: 'most'

Textbook

EYEBALLING TECHNIQUES TO IDENTIFY RED AND WHITE WINES

USING A NON-LINEAR DECISION BOUNDARY LIKE A NEURAL NETWORK



Curved Decision Boundary for Acidity and Total Sulfur Dioxide Related to Wine Color

CONFUSION MATRIX

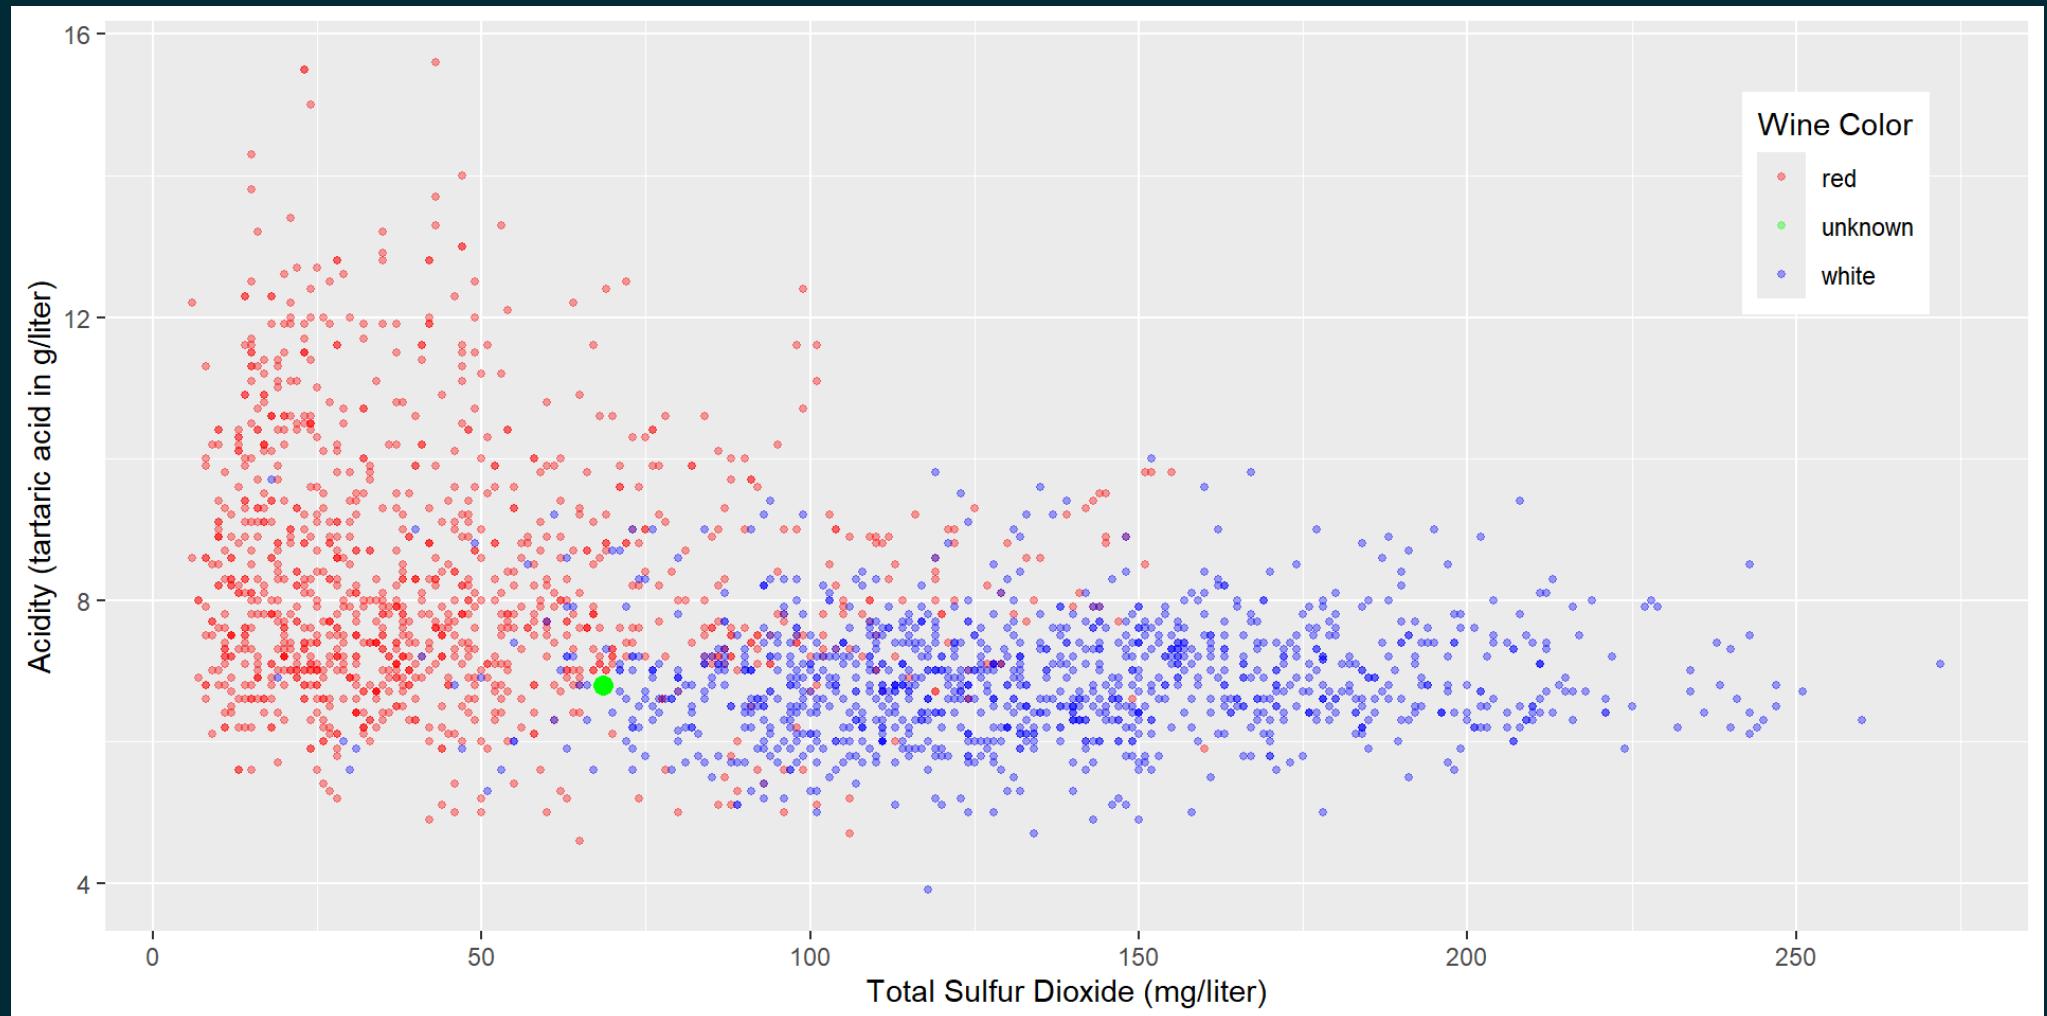
Truth			
Prediction	Red Wine	White Wine	
Red Wine	TP: 'most'	FP: 'few'	
White Wine	FN: 'few'	TN: 'most'	

Textbook

SO, HOW DOES K NEAREST NEIGHBORS WORK?

K NEAREST NEIGHBORS K=1

(FIND THE CLOSEST DATAPoint TO THE ONE WE WANT TO PREDICT)

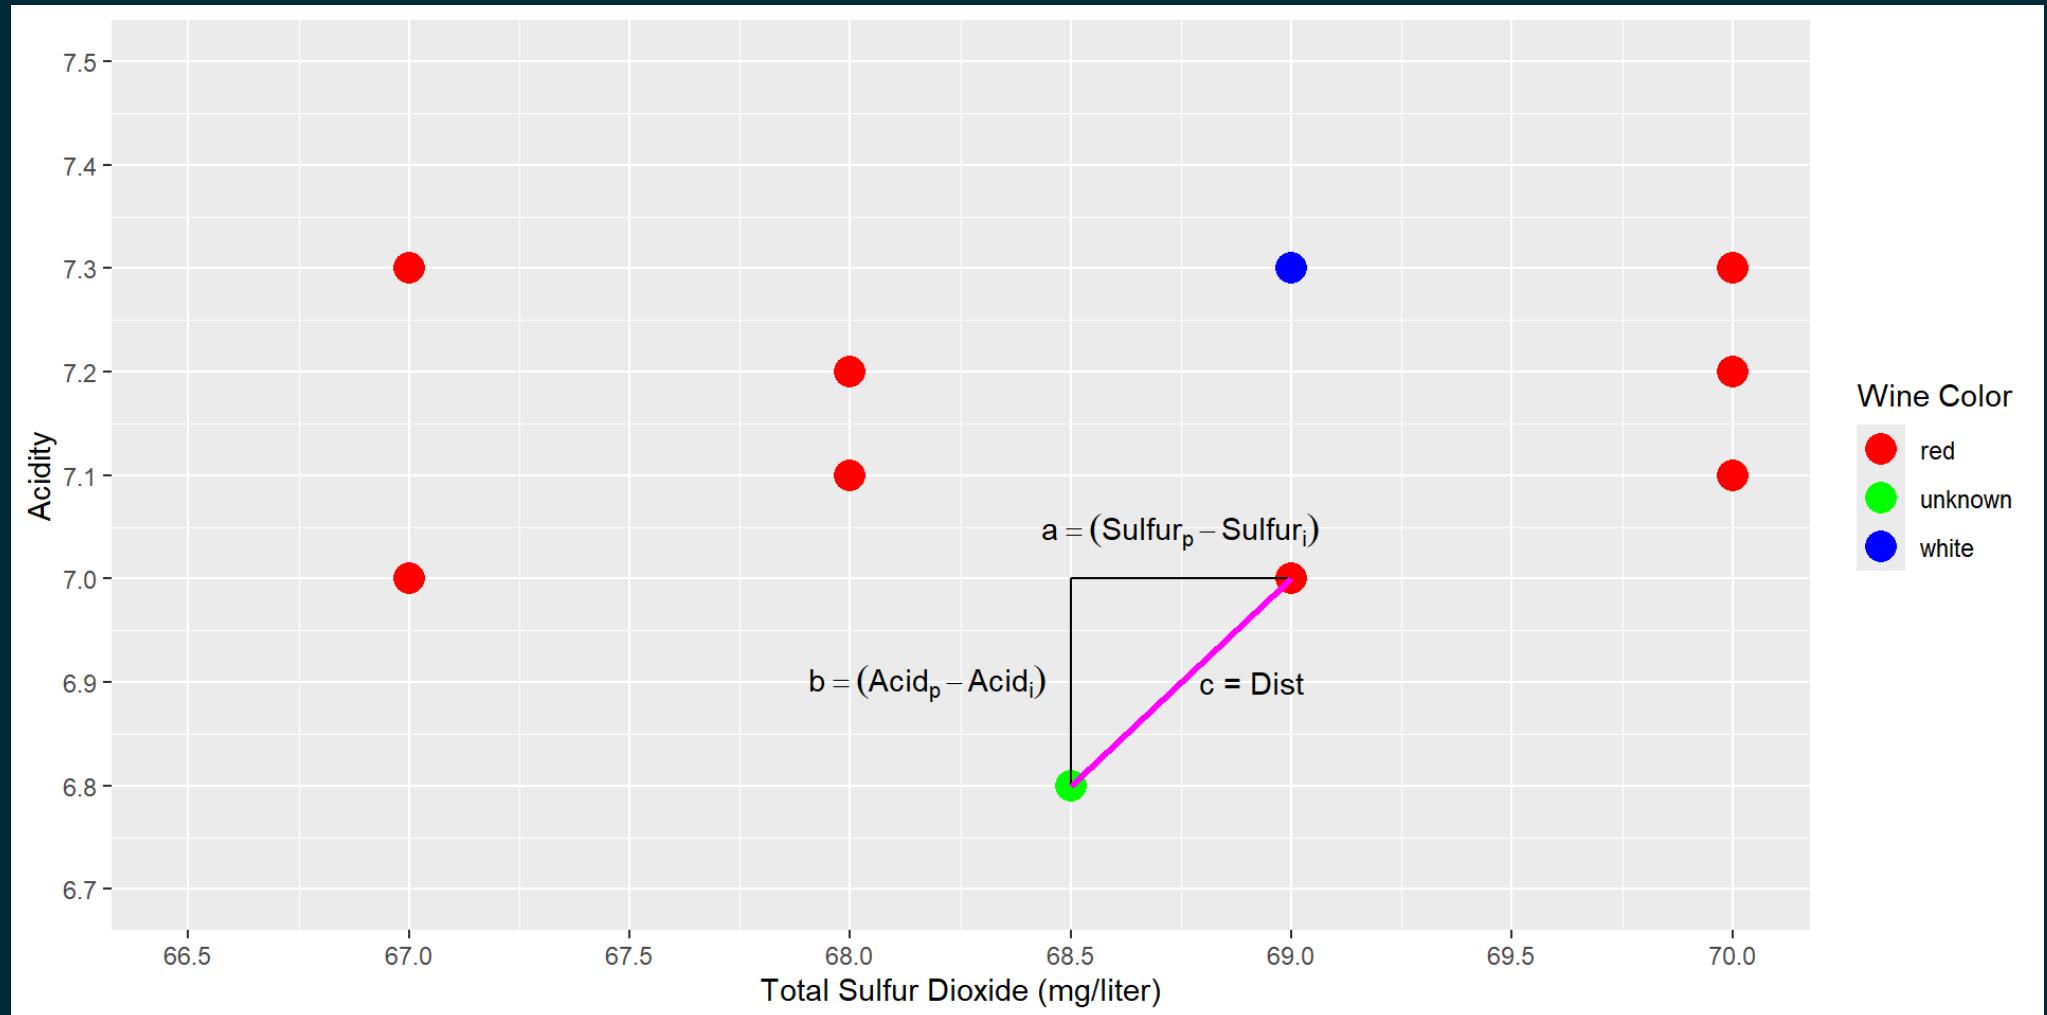


Acidity and Total Sulfur Dioxide Related to Wine Color

Textbook

K NEAREST NEIGHBORS K=1

(FIND THE CLOSEST DATAPoint TO THE ONE WE WANT TO PREDICT)



Predicting Wine Color with k-Nearest Neighbors (k=1)

Textbook

HOW TO CALCULATE EUCLIDEAN DISTANCE FOR TWO VARIABLES

Assume our observations have **two predictor variables** x and y . We compare the unknown point p to one of the points from the training data (e.g., point i):

$$Dist_i = \sqrt{(x_p - x_i)^2 + (y_p - y_i)^2}$$

HOW TO CALCULATE EUCLIDEAN DISTANCE FOR THREE VARIABLES

Assume our observations have **three predictor variables** x , y , and z . We compare the unknown point p to one of the points from the training data (e.g., point i):

$$Dist_i = \sqrt{(x_p - x_i)^2 + (y_p - y_i)^2 + (z_p - z_i)^2}$$

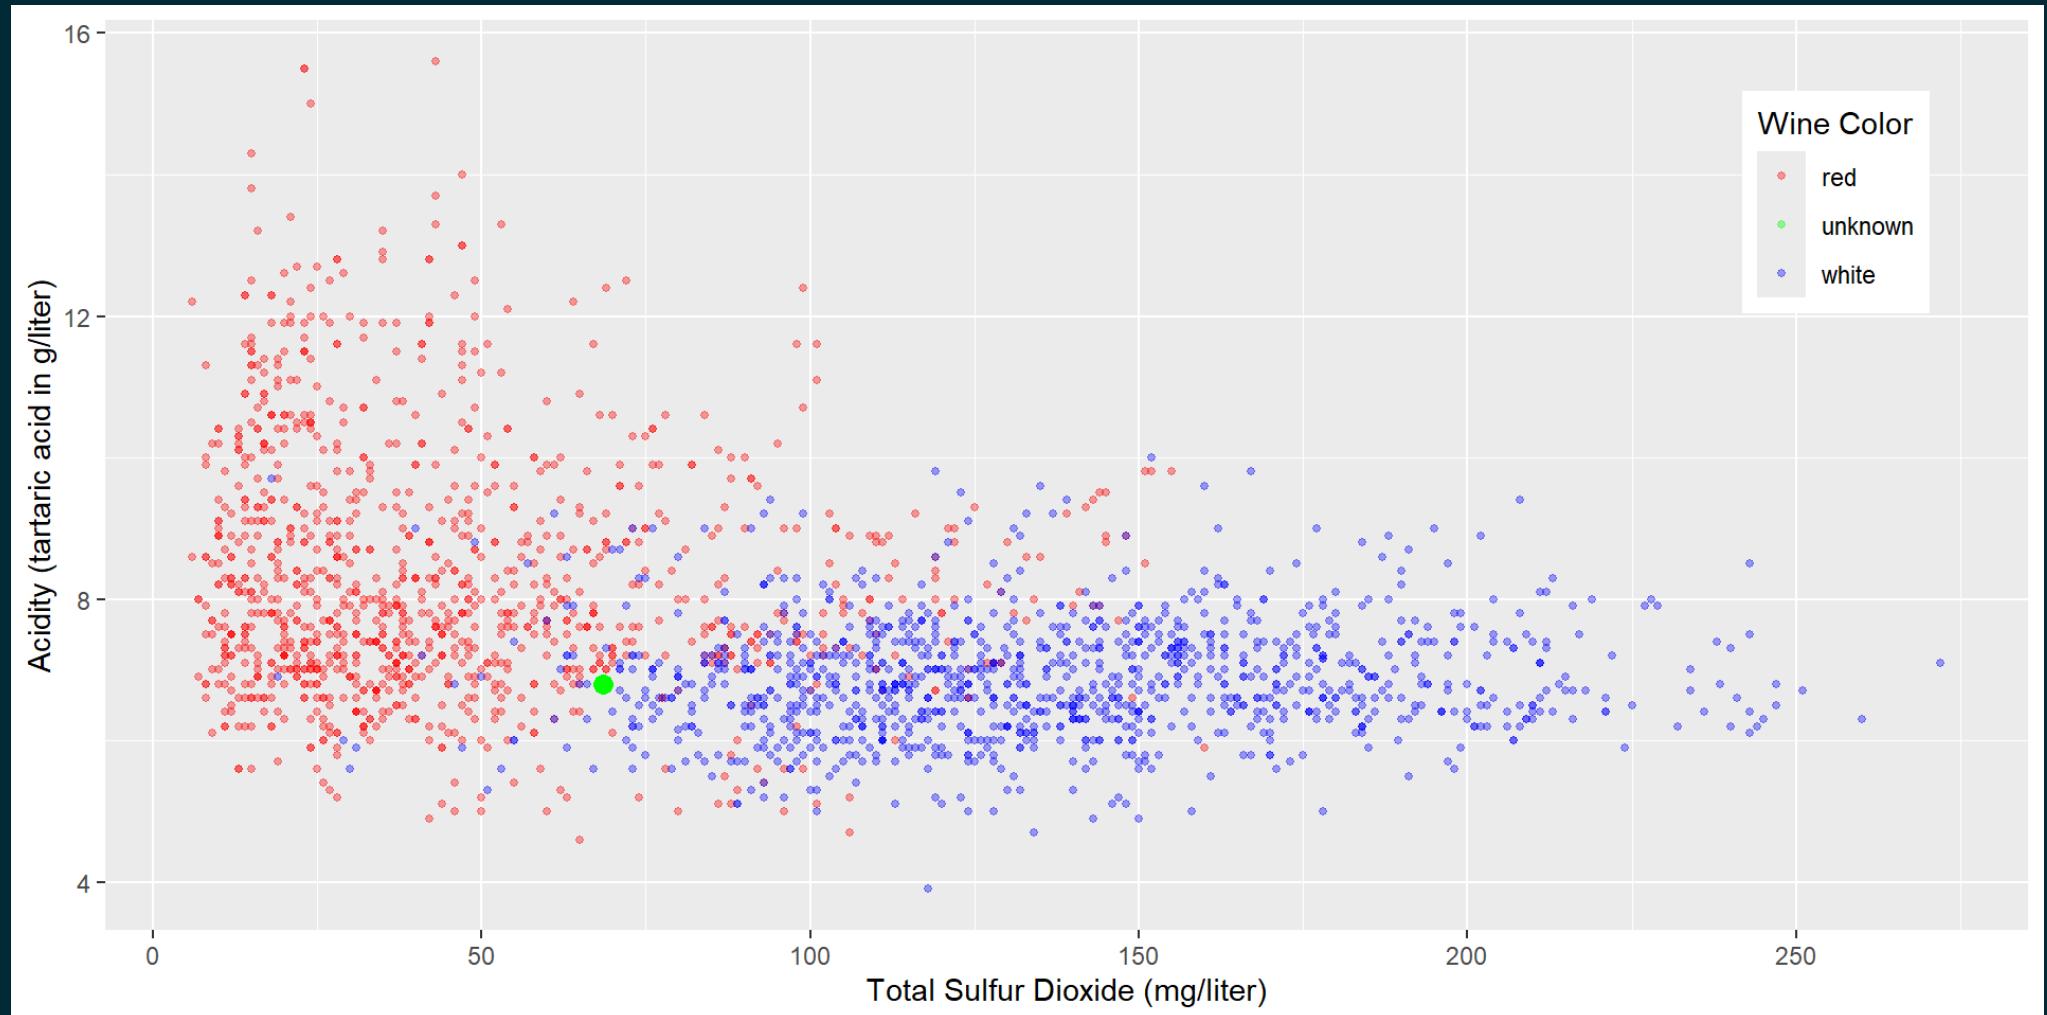
HOW TO CALCULATE EUCLIDEAN DISTANCE FOR N VARIABLES

Assume our observations have N predictor variables v_j with $j = 1 \dots N$. We compare the unknown point p to one of the points from the training data (e.g., point i):

$$Dist_i = \sqrt{\sum_{j=1}^N (v_{p,j} - v_{i,j})^2}$$

K NEAREST NEIGHBORS K=4 (FOR A DIFFERENT UNKNOWN WINE)

(FIND THE CLOSEST 4 DATAPPOINT TO THE ONE WE WANT TO PREDICT)



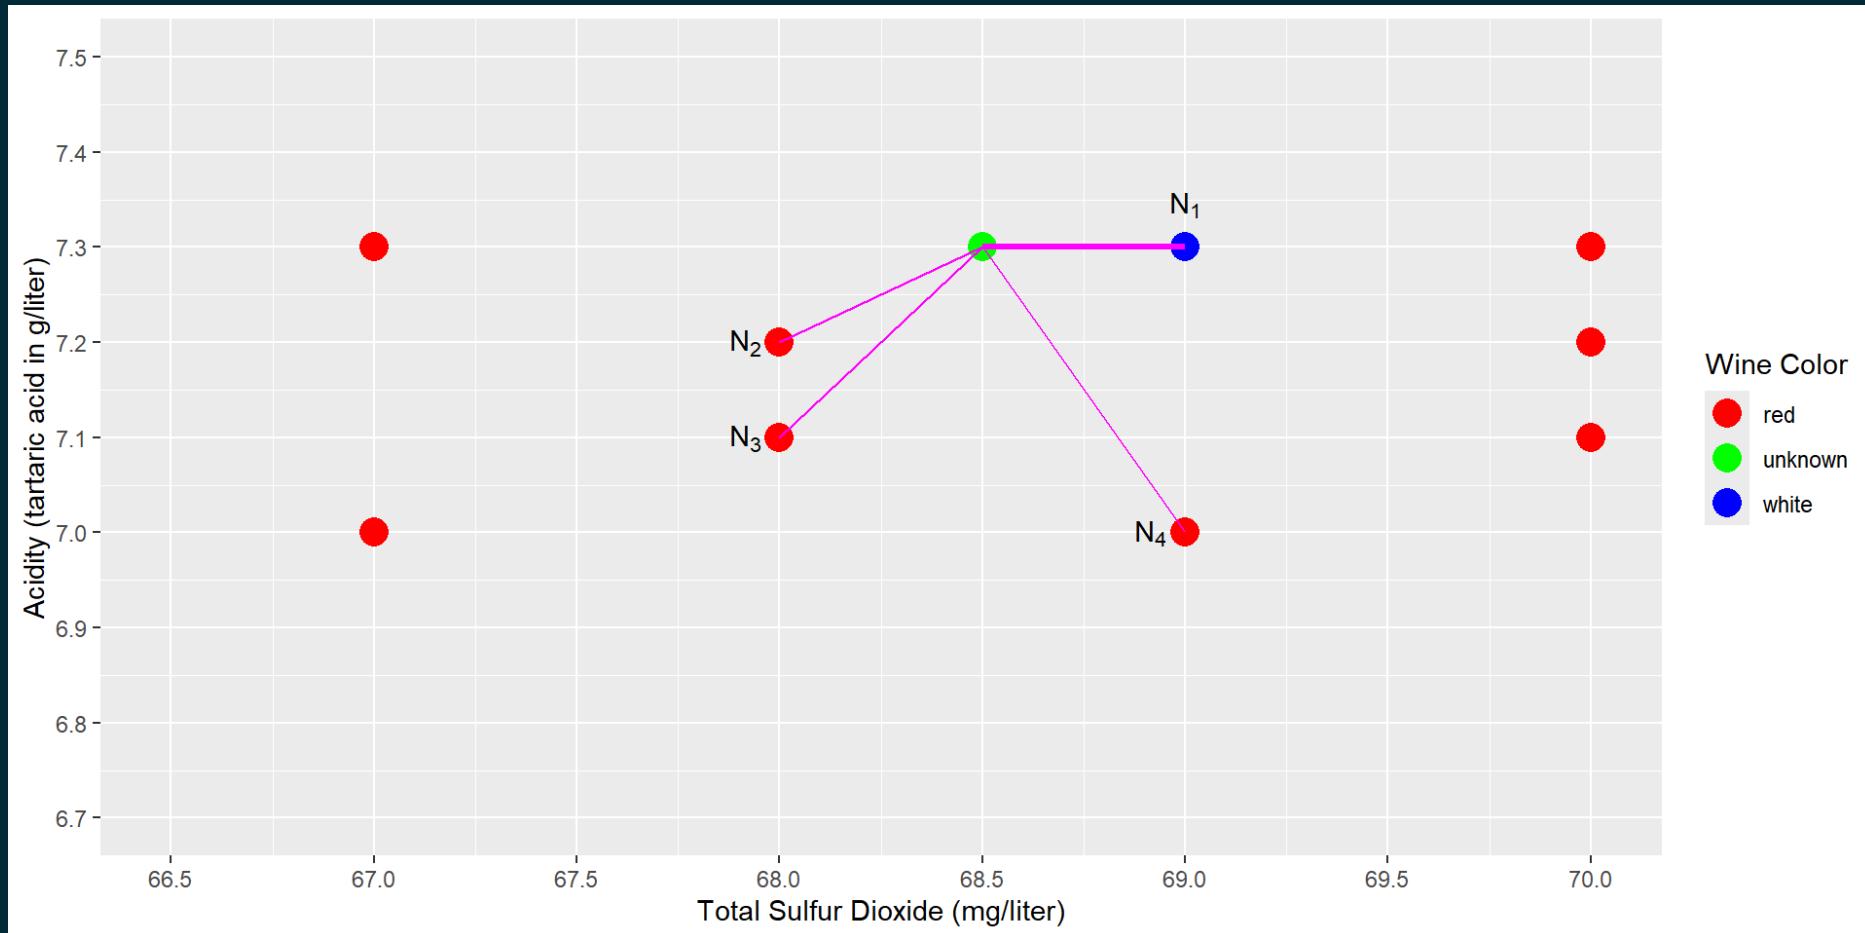
Acidity and Total Sulfur Dioxide Related to Wine Color

Textbook

K NEAREST NEIGHBORS K=4 (FOR A DIFFERENT UNKNOWN WINE)

(FIND THE CLOSEST 4 DATAPPOINT TO THE ONE WE WANT TO PREDICT)

4 NEAREST NEIGHBORS VOTE ON “RED” VS. “WHITE”

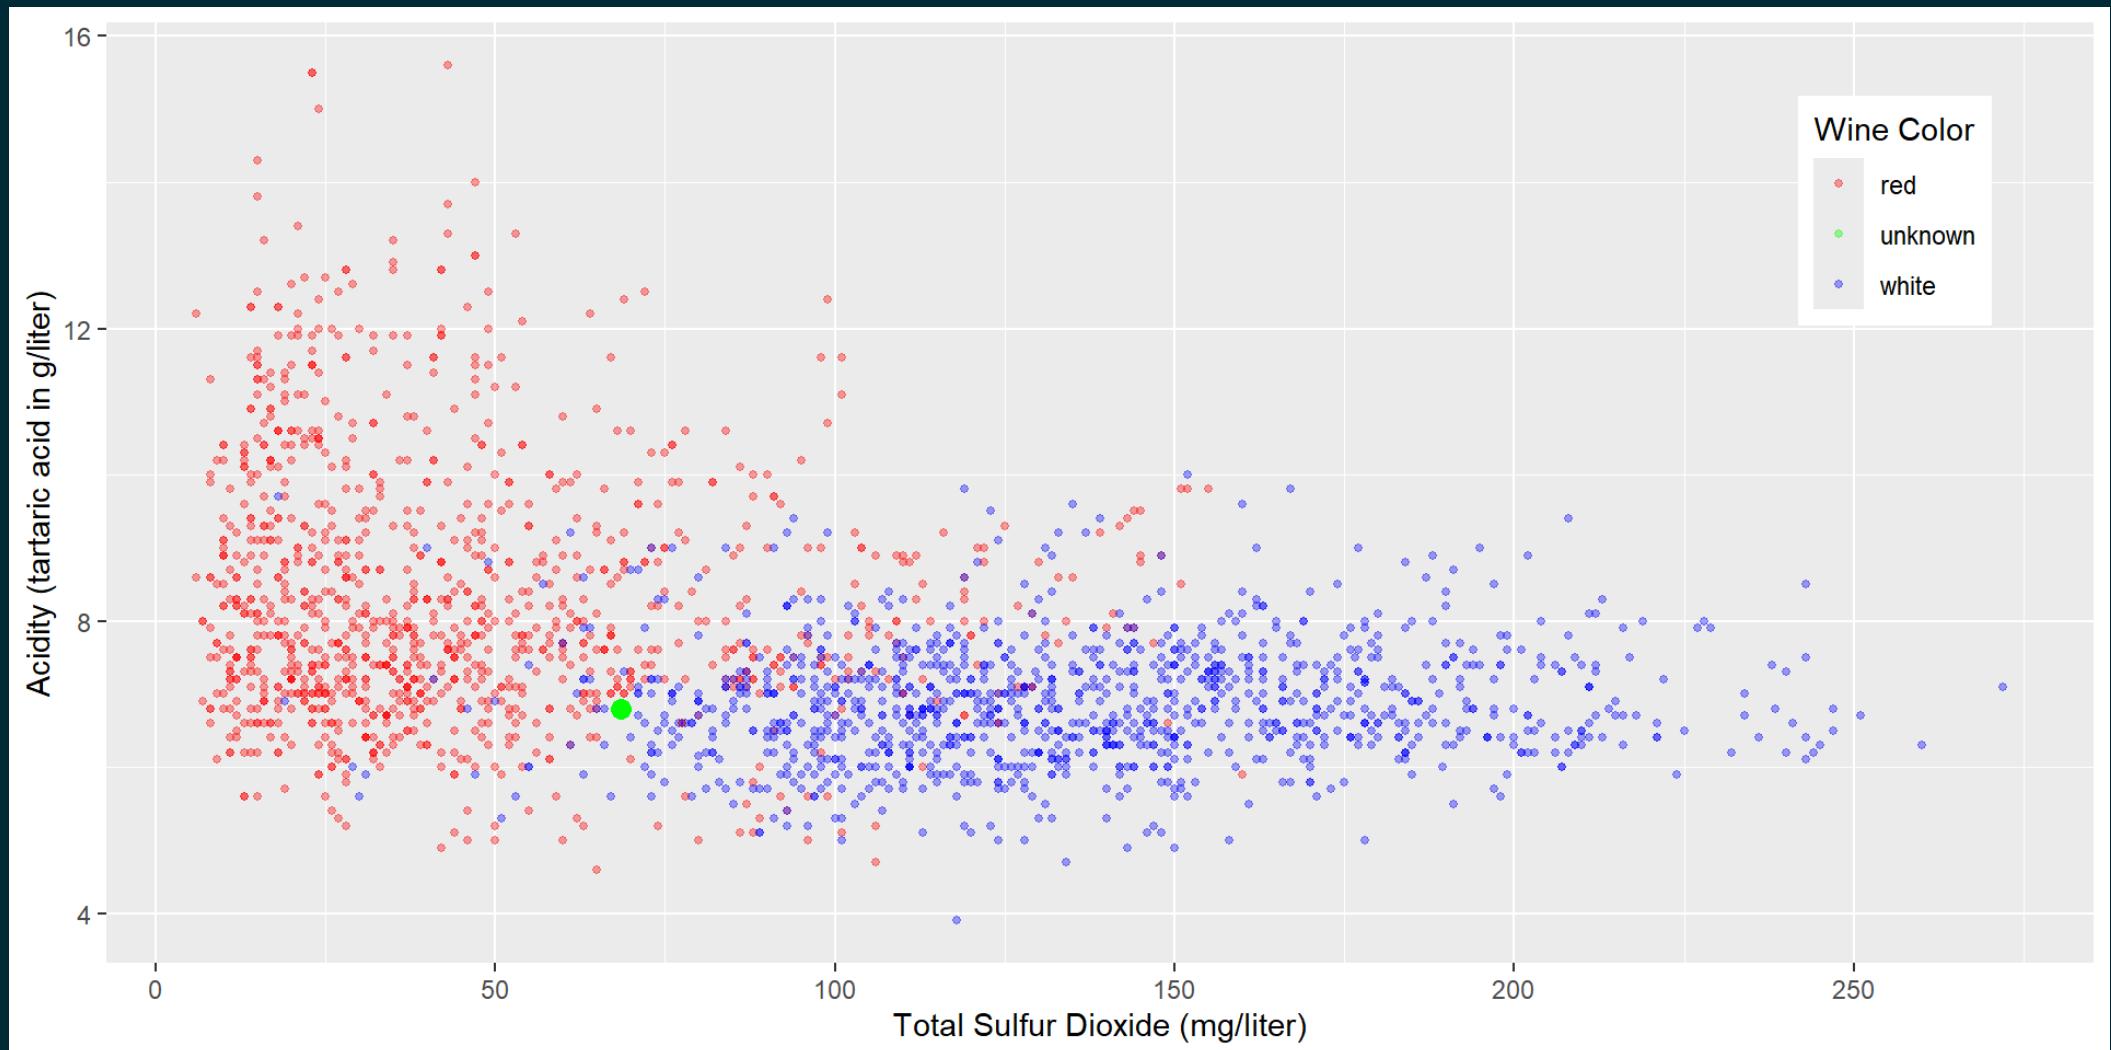


Predicting Wine Color with k-Nearest Neighbors (k=4)

Textbook

K NEAREST NEIGHBORS K=4 (FOR A DIFFERENT UNKNOWN WINE)

WATCH THE SCALE: G/LITER VS. MG/LITER. THAT DOES NOT LOOK RIGHT!



Acidity and Total Sulfur Dioxide Related to Wine Color

Textbook

A FEW COMMON SCALING OPTIONS

- **Same units**

Divide or multiply to get the same units. This is often not possible (e.g., Height and Weight). Or it is not feasible (e.g. Alcohol and StrawberryJuice content in spiked strawberry drink))

- **Rescaling**

Generates a variable y that is scaled to a range between 0 and 1 based on the original variable's value x , its minimum x_{min} and its maximum x_{max} :

$$y = \frac{x - x_{min}}{x_{max} - x_{min}}$$

- **Z-Score Normalization**

Z-score normalization uses the mean (\bar{x}) and the standard deviation (s) of a variable to scale the variable x to the variable z :

$$z = \frac{x - \bar{x}}{s}$$

TIME TO RUN K-NEAREST NEIGHBORS

The required library `tidymodels`, `kknn` and `janitor` are loaded in the background, and the Data we are using are already saved in the data frame `DataWine` (see below):

R Code  Start Over

 Run Code

```
1 library(kableExtra)
2 DataWine=readRDS(url("https://ai.lange-analytics.com/data/WineData.rds")) |>
3     clean_names("upper_camel") |>
4     select(WineColor,Sulfur=TotalSulfurDioxide,Acidity) |>
5     mutate(WineColor=as.factor(WineColor))
6 kable(head(DataWine))
```

GENERATE TRAINING AND TESTING DATA (SPLITTING):

R Code ⟳ Start Over

▷ Run Code

```
1 set.seed(876)
2 Split7030=initial_split(DataWine,prop=0.7,strata = WineColor)
3 DataTrain=training(Split7030)
4 DataTest=testing(Split7030)
5
6 Split7030$in_id[1:20] #only for output; usually not needed
```

DataWine

R Code ⟳ Start Over

▷ Run Code

```
1 head(Split7030$data, n=20)
```

DataTrain

R
Code

⟳ Start
Over

▷ Run
Code

```
1 head(DataTrain)
```

DataTest (not original order)

R
Code

⟳ Start
Over

▷ Run
Code

```
1 head(DataTest)
```

R Code ⟳ Start Over

▷ Run Code

```
1 cat("Proportion of red wines in DataTrain", mean(DataTrain$WineColor=="red"))
2 cat("Proportion of red wines in DataTest", mean(DataTest$WineColor=="red"))
```

DEFINING THE RECIPE

[CLICK HERE TO FIND A REFERENCE LIST FOR VARIOUS Step_ COMMANDS](#)

Recipe: Prepare Data for Analysis:

R Code 



```
1 RecipeWine=recipe(WineColor~Acidity+Sulfur, data = DataTrain) |>
2     step_naomit() |>
3     step_normalize(all_predictors())
```

Or:

R Code 



```
1 RecipeWine=recipe(WineColor~, data = DataTrain) |>
2     step_naomit() |>
3     step_normalize(all_predictors())
```

R Code 



```
1 print(RecipeWine) #red output does not mean error
```

CREATING THE MODEL-DESIGN

[CLICK HERE TO FIND A REFERENCE LIST FOR VARIOUS *MODEL-DESIGNS* COMMANDS](#)

R Code

↻ Start Over

▷ Run Code

```
1 ModelDesignKNN=nearest_neighbor(neighbors = 4, weight_func = "rectangular") |>
  2   set_engine("kknn") |>
  3   set_mode("classification")
  4 print(ModelDesignKNN)
```

`weight_func = "rectangular"` is needed to use a *textbook version* of *k-Nearest-Neighbors*. It can be omitted for research.

ADDING RECIPE & MODEL-DESIGN TO WORKFLOW AND FIT() IT WITH THE TRAINING DATA

Putting it all together in a **fitted workflow**:

R Code 

 Run Code

```
1 WFModelWine=workflow() |>
2     add_recipe(Recipewine) |>
3     add_model(ModelDesignKNN) |>
4     fit(DataTrain)
5 print(WFModelWine)
```

THE PREDICTION PROCEDURE

How to use the **fitted** workflow which contains the training data (`DAtaTrain`) to predict the wine color for the wines in the testing dataset:

1. Start with observation $i = 1$ from `DataTest` (the first observation).
2. Take observation i from `DataTest` and use `Acidity` and `Sulfur` to calculate the *Euclidean distance* to **each** of the observations of `DataTrain`.
3. Isolate the 4 observations with the smallest *Euclidean distance* and use the majority of their wine color as a prediction for observation i from `DataTest` (in case of a par, decide randomly).
4. Increase i by one (i.e., take the next observation from `DataTest`) and go to step 2 (until all `DataTest` observations are processed).

PREDICTING WINE COLOR OF TESTING DATA USING THE predict() COMMAND

Predicting with the fitted workflow using `predict()` (not exactly helpful!):

R Code

⟳ Start Over

▷ Run Code

```
1 predict(WFModelWine, DataTest)
```

PREDICTING WINE COLOR OF TESTING DATA USING THE augment() COMMAND

Predicting with the fitted workflow using `augment()`. The `augment()` command *predicts* and then *augments* `DataTest` with the predictions:

R Code ⟳ Start Over

▷ Run Code

```
1 DataPredWithTestData=augment(WFModelWine, DataTest)
2 head(DataPredWithTestData)
```

HAVING A DATA FRAME WITH `truth` AND `estimate`, WE CAN CALCULATE PERFORMANCE METRICS

Confusion Matrix:

```
R Code      ⌂ Start Over      ▶ Run Code  
1 ConfMatrixWine=conf_mat(DataPredWithTestData, truth = WineColor, estimate = .pred_class)  
2 print(ConfMatrixWine)
```

READING THE CONFUSION MATRIX

		Truth	
Prediction	Red Wine	White Wine	
Red Wine	TP: 436	FP: 46	
White Wine	FN: 44	TN: 434	

- The **positive class** (wine is *predicted red*) is in the **first row** and shows all wines that are *predicted positive* (red). 436 are *predicted correctly* (**TP: True Positives**), and 46 are *predicted incorrectly* as red but are white (**FP: False Positives**).
- The **negative class** (wine is *predicted white*) is in the **second row** and shows all wines that are *predicted negative* (white). 44 are *predicted incorrectly* as white but are red (**FP: False Negatives**), and 434 are *predicted correctly* as white (**TP: True Negatives**).
- Remember:
 - in `tidymodels` *first class is always positive*
 - when determining **TP, FP, FN, TN**, approach matrix from the *prediction side*

Accuracy: Number of wines on diagonal/number of all wines ($\frac{TP+TN}{TP+TN+FP+FN}$):

R Code

↻ Start Over

▷ Run Code

```
1 accuracy(DataPredWithTestData, truth = WineToxicology, estimate = .pred_class)
```

TIDYMODELS MODELING PIPELINE

1. Split the Data into *Training* and *Testing Data*

```
1 set.seed(876)
2 Split7030=initial_split(DataToAnalyzeName, prop=0.7, strata=OutcomeVarName)
3 DataTrain=training(Split7030)
4 DataTest=testing(Split7030)
```

2. Create Recipe

```
1 RecipeName=recipe(OutcomeVarName~Pred1VarName+Pred2VarName, data=DataTrain) |>
2     step_naomit() |>
3     step_Name()
```

3 Create Model-Design

```
1 ModelDesignName=MachLearnModelName() |>
2     set_engine("PackageName") |>
3     set_mode("classification or regression")
```

4. Create Workflow and Fit to Training Data

```
1 WFModelName=workflow() |>
2     add_recipe(RecipeName) |>
3     add_model(ModelDesignName) |>
4     fit(DataTrain)
```

5. Predict with Workflow and Augment Testing Data with Predictions

```
1 DataTestWithPred=augment(WFModelName, DataTest)
```

6. Assess Predictive Quality with Metrics

```
1 MetricsCommand(DataTestWithPred, truth=OutcomeVarName,  
2                  estimate=.pred_class or estimate=.pred)
```

WARNING: BE CAREFUL WITH THE ACCURACY RATE

THE STORY OF DR. NEBULOUS'S GAMBLERS SYSTEM

Dr. Nebulous offers a **97% Machine Learning Gambling Prediction**. Here is how it works: Gamblers can buy a prediction for a fee of \$5. Dr. Nebulous will then run his famous machine learning model and send a closed envelope with the prediction. The gambler is supposed to open the envelope in the casino, right before placing a bet of \$100 on a number in roulette. The envelope contains a message that states either “You will win” or “You will lose”, which allows the gambler to act accordingly by either bet or not bet.

Dr. Nebulous claims that a “clinical trial” of 1000 volunteers, who opened the envelope after they had bet on a number in roulette, shows an accuracy of 97.3%.

How could Dr. Nebulous have such a precise model?

WARNING: BE CAREFUL WITH THE ACCURACY RATE

THE STORY OF DR. NEBULOUS'S GAMBLERS SYSTEM

The trick is Dr. Nebulous's machine learning model uses the *naive prognosis*: It always predicts "You will lose".

Here is the confusion matrix from the 1,000 volunteers trial:

		Truth	
Prediction	Win	Lose	
	Win	0	0
Lose	27	973	

Roulette has 37 numbers to bet on. Chance to win is: $\frac{1}{37} = 0.027$.

Out of the 1000 volunteers, 27 are expected to win $\frac{1}{37}/cdot1000 = 27$, and the others (1000-27=973) are expected to lose.

$$\text{Accuracy} = \frac{0 + 973}{1000} = 0.973$$

WARNING: BE CAREFUL WITH THE ACCURACY RATE

THE STORY OF DR. NEBULOUS'S GAMBLERS SYSTEM

		Truth	
Prediction	Win	Lose	
Win	0	0	
Lose	27	973	

However, when we look at the correct positive and the correct negative rate separately, we see that Dr. Nebulous' accuracy rate (although correct) makes little sense.

- The correct negative rate (**specificity**) is 100%
- The correct positive rate (**sensitivity**) is zero (out of the 27 winners, all were falsely predicted as “You will lose”).

This example shows: When interpreting the confusion matrix, you must look at accuracy, sensitivity, and specificity simultaneously

USING ACCURACY(), SENSITIVITY() AND SPECIFICITY() TO ASSES PREDICTION QUALITY

accuracy()

R Code

⟳ Start Over

▷ Run Code

```
1 conf_mat(DataPredWithTestData, truth = WineColor, estimate = .pred_class)
2 accuracy(DataPredWithTestData, truth = WineColor, estimate = .pred_class)
```

$$Accuracy = \frac{TP+TN}{TP+TN+FP+FN}$$

- Remember:
 - TP and TN are the elements on the main diagonal of the confusion matrix

USING ACCURACY(), SENSITIVITY() AND SPECIFICITY() TO ASSES PREDICTION QUALITY

sensitivity()

R Code

↻ Start Over

▷ Run Code

```
1 conf_mat(DataPredWithTestData, truth = WineColor, estimate = .pred_class)
2 sensitivity(DataPredWithTestData, truth = WineColor, estimate = .pred_class)
```

$$Sensitivity = \frac{TP}{TP+FN} \text{ (first column)}$$

- Remember:
 - in `tidymodels` *first class is always positive*, which is relevant for *Sensitivity*
 - when determining *Sensitivity* and *Specificity* approach matrix from the *Truth* side

USING ACCURACY(), SENSITIVITY() AND SPECIFICITY() TO ASSES PREDICTION QUALITY

specificity()

R Code ⟳ Start Over

▷ Run Code

```
1 conf_mat(DataPredWithTestData, truth = WineColor, estimate = .pred_class)
2 specificity(DataPredWithTestData, truth = WineColor, estimate = .pred_class)
```

$$\text{Specificity} = \frac{TN}{TN+FP} \text{ (second column)}$$

- Remember:
 - in `tidymodels` *second class* is always *negative*, which is relevant for *Specificity*
 - when determining *Sensitivity* and *Specificity* approach matrix from the *Truth* side

USING ALL VARIABLES IN THE WINE DATASET TO PREDICT WINE COLOR

Can we improve by using all predictors?

Try the Interactive Exercise: Chapter 4 – k-Nearest Neighbors: Predicting Wine Color

PROJECT: DESIGN A MACHINE LEARNING WORKFLOW FOR OPTICAL CHARACTER RECOGNITION

You will develop a machine learning model based on *k-Nearest Neighbors* to recognize handwritten digits from images.

You will use the MNIST dataset, a standard dataset for image recognition in machine learning (60,000 images for training and 10,000 images for testing). Developed by LeCun, Cortes, and Burges (2010) based on two datasets from handwritten digits obtained from Census workers and high school students.

[Work with the Interactive Exercise: Chapter 4 – k-Nearest Neighbors Project: Read Images with Handwritten Digits](#)