

K NEAREST NEIGHBORS

BEFORE WE BEGIN: LET US RUN AN EXPERIMENT

I need two volunteers who are willing to reveal their political opinion and for which party they would vote, if we had elections today (ideally one Democrat and one Republican).

The survey questions are below (wait for instructions before you take the survey and record your answers from the digital survey **also on a piece of paper**):

- Survey for all students (except the two volunteers):
<https://tinyurl.com/KNearNeighTrain>
- Survey for the two volunteers
<https://tinyurl.com/KNearNeighTest>

IDEAS ON HOW TO IDENTIFY THE 2 VOLUNTEERS POLITICAL AFFILIATION?

- CL: Ask volunteers what their answers were and find nearest neighbor(s) (volunteers: do not reveal political affiliation yet).
- What were the shortcomings of the survey
- CL: Run R Script

OVERVIEW

In this session you will learn:

1. What is the underlying **idea of k-Nearest Neighbors**
2. How similarity can be measured with **Euclidean distance**
3. Why **scaling predictor variables** is important for some machine learning models
4. Why the **tidymodels package** makes it easy to work with machine learning models
5. How you can define a **recipe** to pre-process data with the **tidymodels** package
6. How you can define a **model-design** with the **tidymodels** package
7. How you can create a machine learning **workflow** with the **tidymodels** package
8. How **metrics** derived from a **confusion matrix** can be used to assess prediction quality
9. Why you have to be careful when interpreting *accuracy*, when you work with **unbalanced observations**
0. How a machine learning model can **process images** and how OCR (Optical Character Recognition) works»

ABOUT THE WINE DATASET

We will work with a publicly available ' wine dataset¹ containing 3,198 observations about different wines and their chemical properties.

Our goal is to develop a k-Nearest Neighbors model that can predict if a wine is red or white based on the wine's chemical properties.»

RAW OBSERVATIONS FROM WINE DATASET

```
1 library(rio)
2 DataWine=import("https://ai.lange-analytics.com/data/WineData.rds")
3 print(DataWine)
```

| | wineColor | acidity | volatileAcidity | citricAcid | residualSugar | Chlorides |
|----|-----------|---------|-----------------|------------|---------------|-----------|
| 1 | red | 10.80 | 0.320 | 0.44 | 1.60 | 0.063 |
| 2 | white | 6.40 | 0.310 | 0.39 | 7.50 | 0.040 |
| 3 | white | 9.40 | 0.280 | 0.30 | 1.60 | 0.045 |
| 4 | white | 8.20 | 0.220 | 0.36 | 6.80 | 0.034 |
| 5 | white | 6.40 | 0.290 | 0.44 | 3.60 | 0.197 |
| 6 | red | 6.70 | 0.855 | 0.02 | 1.90 | 0.064 |
| 7 | red | 11.80 | 0.380 | 0.55 | 2.10 | 0.071 |
| 8 | white | 6.70 | 0.250 | 0.23 | 7.20 | 0.038 |
| 9 | red | 7.50 | 0.380 | 0.57 | 2.30 | 0.106 |
| 10 | red | 7.10 | 0.270 | 0.60 | 2.10 | 0.074 |
| 11 | white | 6.40 | 0.270 | 0.19 | 1.90 | 0.085 |
| 12 | red | 7.80 | 0.600 | 0.26 | 2.00 | 0.080 |
| 13 | red | 8.00 | 0.580 | 0.28 | 3.20 | 0.066 |
| 14 | white | 7.00 | 0.360 | 0.35 | 2.50 | 0.048 |
| 15 | red | 9.90 | 0.440 | 0.46 | 2.20 | 0.091 |
| 16 | white | 7.80 | 0.280 | 0.31 | 2.10 | 0.046 |
| 17 | red | 7.60 | 0.400 | 0.29 | 1.90 | 0.078 |

»

OBSERVATIONS FROM WINE DATASET FOR SELECTED VARIABLES

SULFOR DIOXIDE AND ACIDITY

Note we use `clean_names("upper_camel")` from the `janitor` package to change all column (variable) names to UpperCamel.

```
1 library(tidyverse); library(rio);library(janitor)
2 DataWine=import("https://ai.lange-analytics.com/data/WineData.rds") %>%
3   clean_names("upper_camel") %>%
4   select(WineColor,Sulfur=TotalSulfurDioxide,Acidity) %>%
5   mutate(WineColor=as.factor(WineColor))
6 print(DataWine)
```


| | WineColor | Sulfur | Acidity |
|----|-----------|--------|---------|
| 1 | red | 37.0 | 10.80 |
| 2 | white | 213.0 | 6.40 |
| 3 | white | 139.0 | 9.40 |
| 4 | white | 90.0 | 8.20 |
| 5 | white | 183.0 | 6.40 |
| 6 | red | 38.0 | 6.70 |
| 7 | red | 19.0 | 11.80 |
| 8 | white | 220.0 | 6.70 |
| 9 | red | 12.0 | 7.50 |
| 10 | red | 25.0 | 7.10 |
| 11 | white | 196.0 | 6.40 |
| 12 | red | 131.0 | 7.80 |

BEFORE STARTING WITH K NEAREST NEIGHBORS

LET US FIND SOME EYEBALLING TECHNIQUES THAT ARE RELATED TO VARIOUS MACHINE LEARNING MODELS»

EYE BALLING TECHNIQUES TO IDENTIFY RED AND WHITE WINES

TRY EYEBALLING THE TRAINING DATA ($N = 2238$)




Acidity and Total Sulfur Dioxide Related to Wine Color

<https://ai.lange-analytics.com/>

EYE BALLING TECHNIQUES TO IDENTIFY RED AND WHITE WINES

HORIZONTAL BOUNDARY



Horizontal Decision Boundary for Acidity and Total Sulfur Dioxide Related to Wine Color


CONFUSION MATRIX

| Prediction | Truth | |
|------------|-------|-------|
| | red | white |
| red | 510 | 80 |
| white | 609 | 1039 |

ACCURACY

EYEBALLING TECHNIQUES TO IDENTIFY RED AND WHITE WINES

CREATING SUBSPACES LIKE SIMILAR TO A DECISION TREE



Sub-Space Boundaries for Acidity and Total Sulfur Dioxide Related to Wine Color


CONFUSION MATRIX

| Prediction | Truth | |
|------------|-------|-------|
| | red | white |
| red | 1016 | 145 |
| white | 103 | 974 |

ACCURACY

EYEBALLING TECHNIQUES TO IDENTIFY RED AND WHITE WINES

USING A NON-LINEAR DECISION BOUNDARY LIKE A NEURAL NETWORK




Curved Decision Boundary for Acidity and Total Sulfur Dioxide Related to Wine Color

CONFUSION MATRIX AND ACCURACY

Similar as previous one.


**SO, HOW DOES K NEAREST NEIGHBORS
WORK?**

K NEAREST NEIGHBORS K=1



Acidity and Total Sulfur Dioxide Related to Wine Color»

K NEAREST NEIGHBORS K=1



Predicting Wine Color with k-Nearest Neighbors (k=1)

HOW TO CALCULATE EUCLIDEAN DISTANCE FOR TWO VARIABLES

Assume our observations have **two predictor variables** x and y . We compare the unknown point p to one of the points from the training data (e.g., point i):

$$Dist_i = \sqrt{(x_p - x_i)^2 + (y_p - y_i)^2}$$

»

HOW TO CALCULATE EUCLIDEAN DISTANCE FOR THREE VARIABLES

Assume our observations have **three predictor variables** x , y , and z . We compare the unknown point p to one of the points from the training data (e.g., point i):

$$Dist_i = \sqrt{(x_p - x_i)^2 + (y_p - y_i)^2 + (z_p - z_i)^2}$$

»


HOW TO CALCULATE EUCLIDEAN DISTANCE FOR N VARIABLES

Assume our observations have N predictor variables v_j with $j = 1 \dots N$. We compare the unknown point p to one of the points from the training data (e.g., point i):

$$Dist_i = \sqrt{\sum_{j=1}^N (v_{p,j} - v_{i,j})^2}$$

»

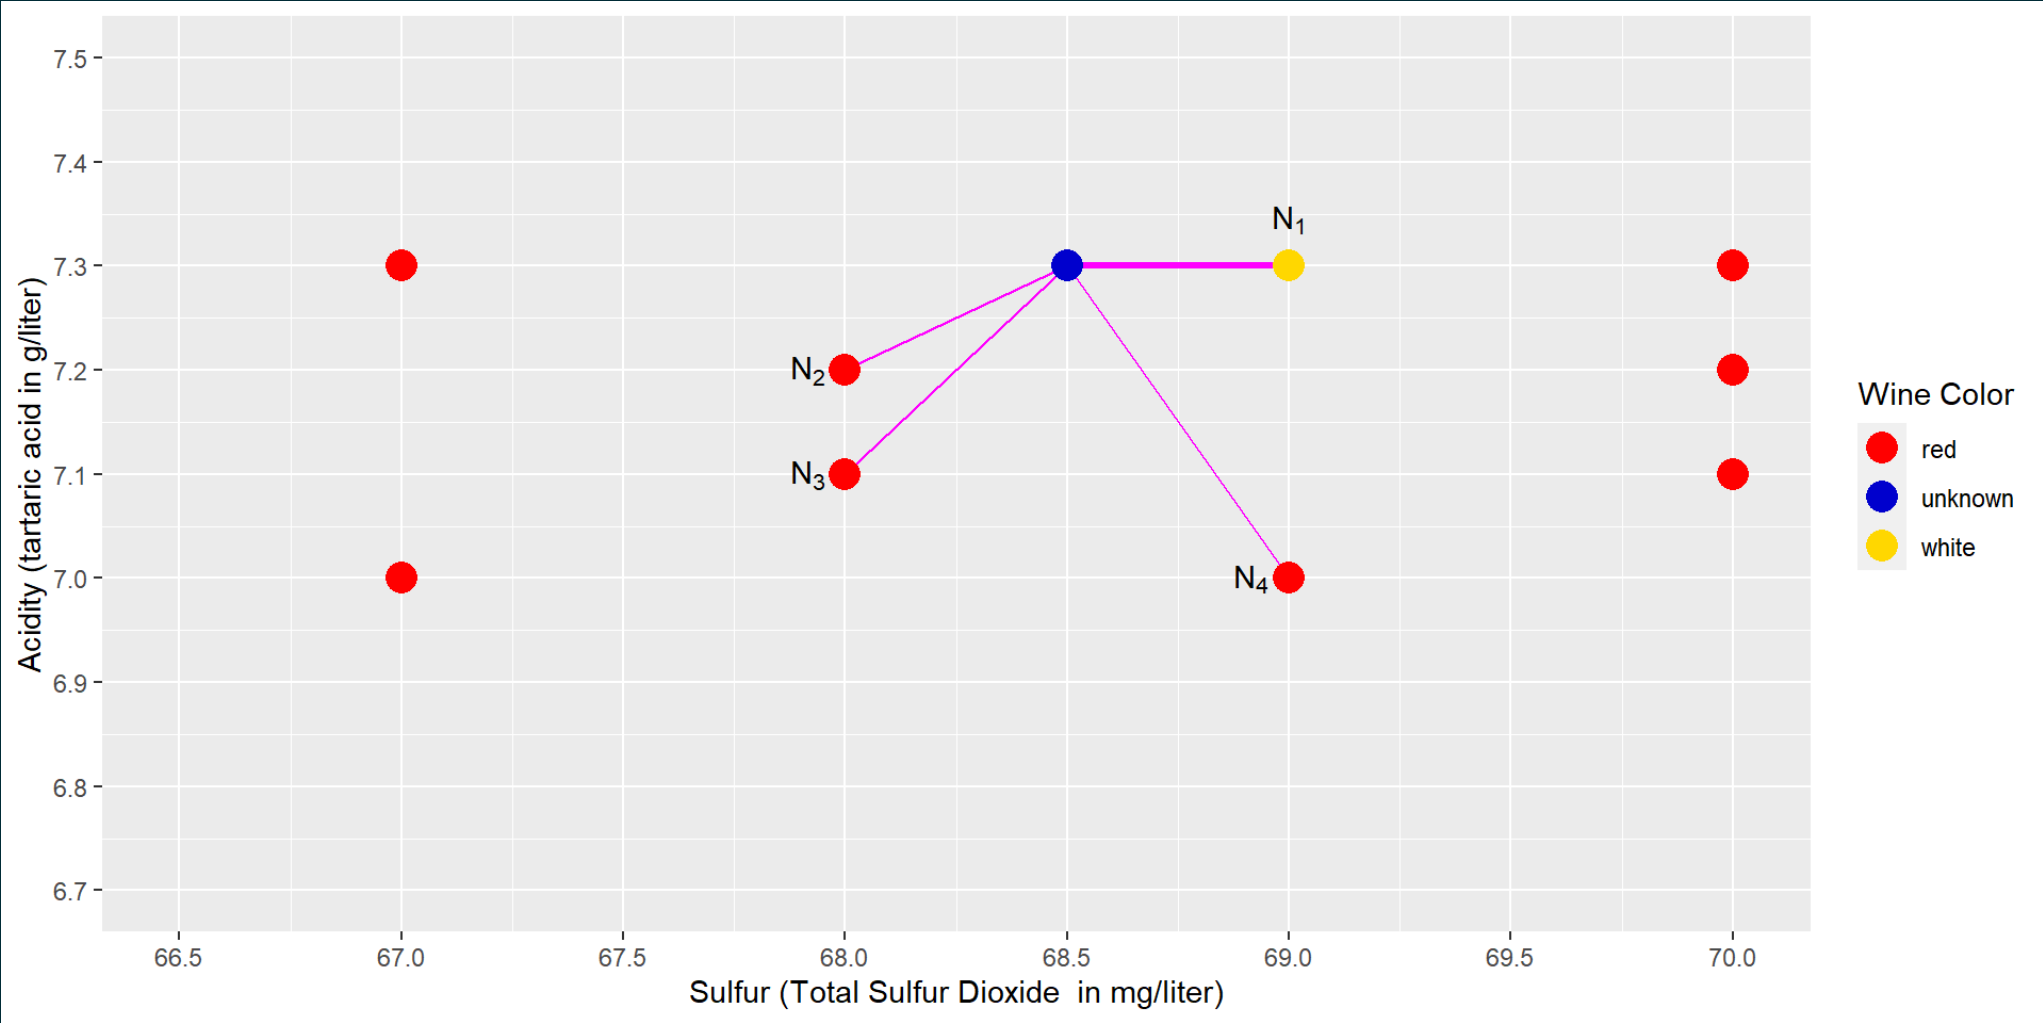
K NEAREST NEIGHBORS $K=4$ (FOR A DIFFERENT UNKNOWN WINE)



Acidity and Total Sulfur Dioxide Related to Wine Color

K NEAREST NEIGHBORS K=4 (FOR A DIFFERENT UNKNOWN WINE)

4 NEAREST NEIGHBORS VOTE ON “RED” VS. “WHITE”




Predicting Wine Color with k-Nearest Neighbors (k=4)

<https://ai.lange-analytics.com/>

K NEAREST NEIGHBORS K=4 (FOR A DIFFERENT UNKNOWN WINE)

WATCH THE SCALE: G/LITER VS. MG/LITER. THAT DOES NOT LOOK RIGHT!



Acidity and Total Sulfur Dioxide Related to Wine Color »

<https://ai.lange-analytics.com/>

A FEW COMMON SCALING OPTIONS

- **Convert from sub-units to major units or vice versa by multiplying or deviding**

For example, if one variable is in feet and one in inches, you can multiply feet by 12 to get inches or devide inches by 12 to get feet. This is often not possible (e.g., **Height** and **Weight**).

- **Rescaling** (R command **range()**)

Generates a variable y that is scaled to a range between 0 and 1 based on the original variable's value x , its minimum x_{min} and its maximum x_{max} :

$$y = \frac{x - x_{min}}{x_{max} - x_{min}}$$

- **Z-Score Normalization**

Z-score normalization uses the mean (\bar{x}) and the standard deviation (s) of a variable to scale the variable x to the variable z :

$$z = \frac{x - \bar{x}}{s}$$

»

TIME TO RUN K-NEAREST NEIGHBORS

LOADING DATA AND SELECTING VARIABLES

```
1 library(tidyverse); library(rio);library(janitor)
2 DataWine=import("https://ai.lange-analytics.com/data/WineData.rds") %>%
3   clean_names("upper_camel") %>%
4   select(WineColor,Sulfur=TotalSulfurDioxide,Acidity) %>%
5   mutate(WineColor=as.factor(WineColor))
6 print(DataWine)
```

| | WineColor | Sulfur | Acidity |
|----|-----------|--------|---------|
| 1 | red | 37.0 | 10.80 |
| 2 | white | 213.0 | 6.40 |
| 3 | white | 139.0 | 9.40 |
| 4 | white | 90.0 | 8.20 |
| 5 | white | 183.0 | 6.40 |
| 6 | red | 38.0 | 6.70 |
| 7 | red | 19.0 | 11.80 |
| 8 | white | 220.0 | 6.70 |
| 9 | red | 12.0 | 7.50 |
| 10 | red | 25.0 | 7.10 |
| 11 | white | 196.0 | 6.40 |
| 12 | red | 131.0 | 7.80 |
| 13 | red | 114.0 | 8.00 |
| 14 | white | 161.0 | 7.00 |
| 15 | red | 41.0 | 9.90 |
| 16 | white | 208.0 | 7.80 |
| 17 | red | 66.0 | 7.60 |

TIME TO RUN K-NEAREST NEIGHBORS

Generate Training and Testing Data (Splitting):

```
1 library(tidymodels);
2 set.seed(876)
3 Split7030=initial_split(DataWine,prop=0.7,strata = WineColor)
4
5 DataTrain=training(Split7030)
6 DataTest=testing(Split7030)
7 print(DataTrain)
```

| | WineColor | Sulfur | Acidity |
|----|-----------|--------|---------|
| 1 | red | 37.0 | 10.80 |
| 2 | red | 38.0 | 6.70 |
| 3 | red | 12.0 | 7.50 |
| 4 | red | 25.0 | 7.10 |
| 5 | red | 114.0 | 8.00 |
| 6 | red | 66.0 | 7.60 |
| 7 | red | 49.0 | 6.80 |
| 8 | red | 110.0 | 7.00 |
| 9 | red | 44.0 | 6.50 |
| 10 | red | 10.0 | 10.40 |
| 11 | red | 149.0 | 6.60 |
| 12 | red | 31.0 | 6.60 |
| 13 | red | 49.0 | 6.60 |
| 14 | red | 17.0 | 9.10 |
| 15 | red | 32.0 | 6.10 |
| 16 | red | 30.0 | 8.50 |
| 17 | red | 17.0 | 10.20 |

```
1 print(DataTest)
```

| | WineColor | Sulfur | Acidity |
|----|-----------|--------|---------|
| 1 | white | 90.0 | 8.20 |
| 2 | red | 19.0 | 11.80 |
| 3 | white | 220.0 | 6.70 |
| 4 | red | 131.0 | 7.80 |
| 5 | white | 161.0 | 7.00 |
| 6 | red | 41.0 | 9.90 |
| 7 | white | 156.0 | 7.80 |
| 8 | white | 150.0 | 6.50 |
| 9 | red | 102.0 | 7.90 |
| 10 | red | 11.0 | 5.80 |
| 11 | white | 64.0 | 5.90 |
| 12 | white | 178.0 | 6.90 |
| 13 | red | 41.0 | 12.60 |
| 14 | red | 19.0 | 12.70 |
| 15 | white | 206.0 | 6.20 |
| 16 | white | 149.0 | 6.90 |
| 17 | white | 182.0 | 7.80 |

TIME TO RUN K-NEAREST NEIGHBORS

CLICK HERE TO FIND A REFERENCE LIST FOR VARIOUS **Step_** COMMANDS

Recipe: Prepare Data for Analysis:

```
1 RecipeWine=recipe(WineColor~Acidity+Sulfur, data = DataTrain) %>%
2     step_naomit() %>%
3     step_normalize(all_predictors())
4 # or
5
6 RecipeWine=recipe(WineColor~., data = DataTrain) %>%
7     step_naomit() %>%
8     step_normalize(all_predictors())
9 print(RecipeWine)
```

TIME TO RUN K-NEAREST NEIGHBORS

CLICK HERE TO FIND A REFERENCE LIST FOR VARIOUS **Step_** COMMANDS

Creating a Model Design:

```
1 ModelDesignKNN=nearest_neighbor(neighbors = 4, weight_func = "rectangular") %>%
2     set_engine("knn") %>%
3     set_mode("classification")
4 print(ModelDesignKNN)
```

K-Nearest Neighbor Model Specification (classification)

Main Arguments:

```
neighbors = 4
weight_func = rectangular
```

Computational engine: knn

TIME TO RUN K-NEAREST NEIGHBORS

Putting it all together in a **fitted workflow**:

```
1 WfModelWine=workflow() %>%
2       add_recipe(RecipeWine) %>%
3       add_model(ModelDesignKNN) %>%
4       fit(DataTrain)
5 print(WfModelWine)
```

```
== Workflow [trained] ==
Preprocessor: Recipe
Model: nearest_neighbor()
```

```
— Preprocessor —
2 Recipe Steps
```

- step_naomit()
- step_normalize()

```
— Model —
```

```
Call:
kkn::train.kkn(formula = ..y ~ ., data = data, ks = min_rows(4, data, 5), kernel =
~"rectangular")
```

```
Type of response variable: nominal
Minimal misclassification: 0.1000894
```

TIME TO RUN K-NEAREST NEIGHBORS

How to use the **fitted workflow** to predict the wine color for the wines in the testing dataset:

1. Start with observation $i = 1$ from **DataTest** (the first observation).
2. Take observation i from **DataTest** and use **Acidity** and **Sulfur** to calculate the Euclidean distance to **each** of the observations of **DataTrain**.
3. Isolate the 4 observations with the smallest Euclidean distance and use the majority of their wine color as a prediction for observation i from **DataTest** (in case of a par, decide randomly).
4. Increase i by one (i.e., take the next observation from **DataTest**) and go to step 2 (until all **DataTest** observations are processed).

TIME TO RUN K-NEAREST NEIGHBORS

Predicting with the fitted workflow using `predict()` (not exactly helpful!):

```
1 predict(WFModelWine, DataTest)
```

```
# A tibble: 960 × 1
```

```
  .pred_class
```

```
  <fct>
```

```
1 white
```

```
2 red
```

```
3 white
```

```
4 white
```

```
5 white
```

```
6 red
```

```
7 white
```

```
8 white
```

```
9 red
```

```
10 red
```

```
# i 950 more rows
```

TIME TO RUN K-NEAREST NEIGHBORS

Predicting with the fitted workflow using `augment()` which *augments* `DataTest` with the predictions:

```
1 DataPredWithTestData=augment(WFModelWine, DataTest)
2 head(DataPredWithTestData)
```

```
# A tibble: 6 × 6
  WineColor Sulfur Acidity .pred_class .pred_red .pred_white
  <fct>      <dbl>   <dbl> <fct>      <dbl>      <dbl>
1 white      90      8.2 white      0.25      0.75
2 red        19     11.8 red         1         0
3 white     220      6.7 white      0         1
4 red       131      7.8 white     0.25     0.75
5 white     161      7   white      0         1
6 red        41      9.9 red         1         0
```


HAVING A DATA FRAME WITH **truth** AND **esimate** WE CAN CALCULATE PERFORMANCE METRICS

Confusion Matrix:

```
1 ConfMatrixWine=conf_mat(DataPredWithTestData, truth = WineColor, estimate = .pred_clas
2 print(ConfMatrixWine)
```

| | Truth | |
|------------|-------|-------|
| Prediction | red | white |
| red | 436 | 46 |
| white | 44 | 434 |

READING THE CONFUSION MATRIX

| | Truth | |
|------------|----------|------------|
| Prediction | Red Wine | White Wine |
| Red Wine | TP: 436 | FP: 46 |
| White Wine | FN: 44 | TN: 434 |

- The **positive class** (wine is “red”) is in the **first column**. 436 of the positives are classified correctly (TR: true positives), and 44 positives are incorrectly classified (FN: false negatives).
- The **negative class** (wine is “white”) is in the **second column**. 44 negatives are incorrectly classified (FP: false positives), and 434 negatives are classified correctly (TN: true negatives).

Accuracy: Number of wines on diagonal/number of all wines:

```
1 accuracy(DataPredWithTestData, truth = WineColor, estimate = .pred_class)

# A tibble: 1 × 3
  .metric .estimator .estimate
  <chr>    <chr>         <dbl>
1 accuracy binary      0.906
```

<https://ai.lange-analytics.com/>

WARNING: BE CAREFUL WITH THE ACCURACY RATE

THE STORY OF DR. NEBULOUS'S GAMBLERS SYSTEM

Dr. Nebulous offers a **97% Machine Learning Gambling Prediction**. Here is how it works: Gamblers can buy a prediction for a fee of \$5. Dr. Nebulous will then run his famous machine learning model and send a closed envelope with the prediction. The gambler is supposed to open the envelope in the casino, right before placing a bet of \$100 on a number in roulette. The envelope contains a message that states either “You will win” or “You will lose”, which allows the gambler to act accordingly by either bet or not bet.

Dr. Nebulous claims that a “clinical trial” of 1000 volunteers, who opened the envelope after they had bet on a number in roulette, shows an accuracy of 97.3%.

How could Dr. Nebulous have such a precise model?

<https://ai.lange-analytics.com/>

WARNING: BE CAREFUL WITH THE ACCURACY RATE

THE STORY OF DR. NEBULOUS'S GAMBLERS SYSTEM

The trick is Dr. Nebulous's machine learning model uses the *naive prognosis*: It always predicts "You will lose".

Here is the confusion matrix from the 1,000 volunteers trial:

| Prediction \ Truth | Win | Lose |
|--------------------|-----|------|
| | Win | Lose |
| Win | 0 | 0 |
| Lose | 27 | 997 |

Roulette has 37 numbers to bet on. Chance to win is: $\frac{1}{37} = 0.027$.

Out of the 1000 volunteers, 27 are expected to win, and 973 are expected to lose.

$$Accuracy = \frac{0 + 973}{1000} = 0.973$$

WARNING: BE CAREFUL WITH THE ACCURACY RATE

THE STORY OF DR. NEBULOUS'S GAMBLERS SYSTEM

| Prediction | Truth | |
|------------|-------|------|
| | Win | Lose |
| Win | 0 | 0 |
| Lose | 27 | 997 |

However, when we look at the correct positive and the correct negative rate separately, we see that Dr. Nebulous' accuracy rate (although correct) makes little sense.

- The correct negative rate (**specificity**) is 100%
- The correct positive rate (**sensitivity**) is zero (out of the 27 winners, all were falsely predicted as “You will lose”).

This example shows: When interpreting the confusion matrix, you must look at accuracy, sensitivity, and specificity simultaneously

TIME TO RUN K-NEAREST NEIGHBORS 🤖

`accuracy()`, `sensitivity()` and `specificity()` for the wine data:

```
1 accuracy(DataPredWithTestData, truth = WineColor, estimate = .pred_class)
```

```
# A tibble: 1 × 3  
  .metric .estimator .estimate  
  <chr>    <chr>         <dbl>  
1 accuracy binary         0.906
```

```
1 sensitivity(DataPredWithTestData, truth = WineColor, estimate = .pred_class)
```

```
# A tibble: 1 × 3  
  .metric .estimator .estimate  
  <chr>    <chr>         <dbl>  
1 sensitivity binary         0.908
```

```
1 specificity(DataPredWithTestData, truth = WineColor, estimate = .pred_class)
```

```
# A tibble: 1 × 3  
  .metric .estimator .estimate  
  <chr>    <chr>         <dbl>  
1 specificity binary         0.904
```

Can we improve by using all predictors.»

PROJECT: DESIGN A MACHINE LEARNING WORKFLOW FOR OPTICAL CHARACTER RECOGNITION »


MNIST DATA SET

You will develop a machine learning model based on *k*-Nearest Neighbors to recognize handwritten digits from images.

You will use the MNIST dataset, a standard dataset for image recognition in machine learning (60,000 images for training and 10,000 images for testing). Developed by LeCun, Cortes, and Burges (2010) based on two datasets from handwritten digits obtained from Census workers and high school students.

We will use only the first 500 images of the original MNIST dataset to speed up the *k*-Nearest Neighbors model's training time.

VISUALIZATION OF THE FIRST SIX IMAGES FROM THE MNIST DATA SET



HOW A IMAGE IS STORED IN THE MNIST DATASET




Image of a Handwritten Nine

The image has 28 rows and 28 columns. Each of the 784 cells (pixels) holds a value between 0 (black) and 255 (white)

HOW A IMAGE IS STORED IN THE MNIST DATASET

[illegible]

- Pixel values for a single image are not stored in a table. Otherwise we would end-up with a table containing tables.
- Pixel values are stored as one row for each image.
- Concatenating the 28 rows of an image into one row with $28*28=784$ cells (pixels)

Image of a Handwritten Nine

THREE ROWS FROM THE DATA FRAME OF THE MNIST DATASET

```
1 print(Mnist4PlotAndTable[1:3,1:784])
```

| | Label | Pix1 | Pix2 | Pix3 | Pix4 | Pix5 | Pix6 | Pix7 | Pix8 | Pix9 | Pix10 | Pix11 | Pix12 | Pix13 |
|---|-------|-------|-------|-------|-------|-------|-------|-------|-------|-------|-------|-------|-------|-------|
| 1 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 |
| 2 | 5 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 |
| 3 | 3 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 |
| | Pix14 | Pix15 | Pix16 | Pix17 | Pix18 | Pix19 | Pix20 | Pix21 | Pix22 | Pix23 | Pix24 | Pix25 | Pix26 | |
| 1 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | |
| 2 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | |
| 3 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | |
| | Pix27 | Pix28 | Pix29 | Pix30 | Pix31 | Pix32 | Pix33 | Pix34 | Pix35 | Pix36 | Pix37 | Pix38 | Pix39 | |
| 1 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | |
| 2 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | |
| 3 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | |
| | Pix40 | Pix41 | Pix42 | Pix43 | Pix44 | Pix45 | Pix46 | Pix47 | Pix48 | Pix49 | Pix50 | Pix51 | Pix52 | |
| 1 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | |
| 2 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | |
| 3 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | |
| | Pix53 | Pix54 | Pix55 | Pix56 | Pix57 | Pix58 | Pix59 | Pix60 | Pix61 | Pix62 | Pix63 | Pix64 | Pix65 | |
| 1 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | |
| 2 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | |
| 3 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | |
| | Pix66 | Pix67 | Pix68 | Pix69 | Pix70 | Pix71 | Pix72 | Pix73 | Pix74 | Pix75 | Pix76 | Pix77 | Pix78 | |
| 1 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | |
| 2 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | |
| 3 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | |
| | Pix79 | Pix80 | Pix81 | Pix82 | Pix83 | Pix84 | Pix85 | Pix86 | Pix87 | Pix88 | Pix89 | Pix90 | Pix91 | |

GO TO PROJECT IN BOOK 

BUILD YOUR OWN OCR SYSTEM.»

